

# Systematic Exploration of Critical Software for Catastrophic Cyber-Physical Malware

Suresh Kothari – [kothari@iastate.edu](mailto:kothari@iastate.edu)  
Benjamin Holland – [bholland@iastate.edu](mailto:bholland@iastate.edu)

**milcom**  
Military Communications for the 21st Century

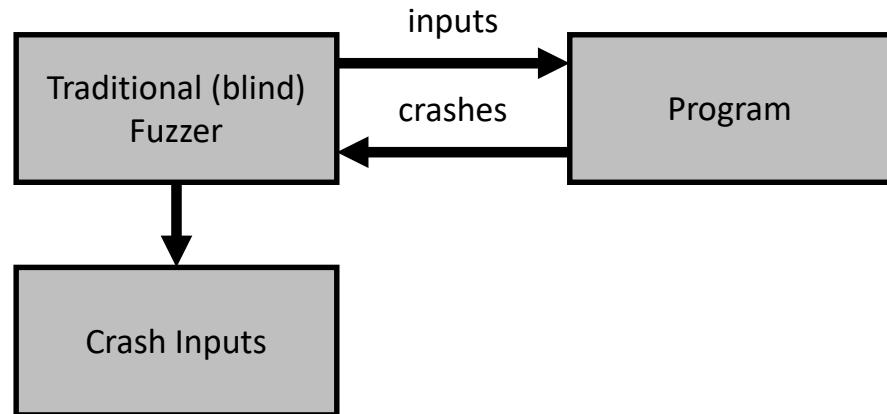
LAX Marriott, Los Angeles, CA

# How do we analyze a program?

- Two main camps
  - Dynamic analysis
    - Run the program with some inputs and see what it does
    - Advantage: Everything we observe is feasible (we just saw it happen)
    - Concern: Input space is HUGE
    - Concern: Did we test the *interesting* inputs?
  - Static analysis
    - Don't run the program, dissect the logic and examine program artifacts
    - Advantage: Bird's eye view of everything that could possibly happen during execution
    - Concern: Number of program behaviors is HUGE
    - Concern: Is it feasible to reach/trigger an artifact of concern?
- What are we looking for?
  - Bugs: Memory corruption, rounding errors, null pointers, infinite loops, stack overflows, race conditions, memory leaks, business logic flaws, ...
  - Not every issue translates to a crash!

# Traditional/Dumb (Blind) Fuzzing

- Start with a test corpus of well formed program inputs
- Apply random or systematic mutations to program inputs
- Run program with mutated inputs and observe whether or not the program crashes
- Repeat until the program crashes
- Input space
  - Reading data in loops could make the input space infinite
  - There are  $2^n$  possible inputs for a binary input of length  $n$



This is about all we can do without examining program artifacts...

# Building a Static Analysis Tool

- Compiler: Lexer → Parser → AST → Semant → Code Generation
- Static Analysis: Lexer → Parser → AST → Semant → Graph of Program Artifacts → Graph Queries of Concerning Program Relationships
- Demo Eclipse Plugin: <http://ben-holland.com/AtlasBrainfuck/updates>

# Brainf\*ck Language

- Designed by Urban Müller in 1992 with the goal of implementing the smallest possible compiler.
- Compiler can be implemented in less than 100 bytes
- Implements a Turing machine

Reference:

<https://en.wikipedia.org/wiki/Brainfuck>

Character	Meaning
>	increment the data pointer (to point to the next cell to the right).
<	decrement the data pointer (to point to the next cell to the left).
+	increment (increase by one) the byte at the data pointer.
-	decrement (decrease by one) the byte at the data pointer.
.	output the byte at the data pointer.
,	accept one byte of input, storing its value in the byte at the data pointer.
[	if the byte at the data pointer is zero, then instead of moving the instruction pointer forward to the next command, jump it <i>forward</i> to the command after the <i>matching</i> ] command.
]	if the byte at the data pointer is nonzero, then instead of moving the instruction pointer forward to the next command, jump it <i>back</i> to the command after the <i>matching</i> [ command.

# Brainf\*ck (Hello World)

```
++++++[>++++[>++>+++>+++>+<<<<-]>+>+>->>+[<]<-]>>. >---. +++++++..+++.>>. <-.<.+++.-----.-  
-----.>>+.>++.
```

# Brainf\*ck Lexical Analysis

```
MOVE_RIGHT: '>';  
MOVE_LEFT: '<';  
INCREMENT: '+';  
DECREMENT: '-';  
WRITE: '.';  
READ: ',';  
LOOP_HEADER: '[';  
LOOP_FOOTER: ']' ;
```

Program: ***++[>+[+]]***.

Program Tokens: INCREMENT INCREMENT LOOP\_HEADER MOVE\_RIGHT  
INCREMENT LOOP\_HEADER INCREMENT LOOP\_FOOTER LOOP\_FOOTER WRITE <EOF>

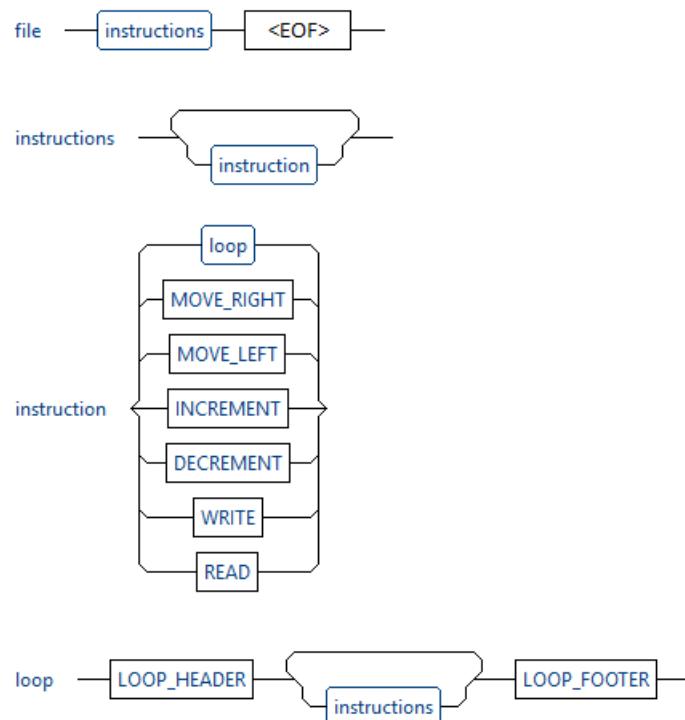
# Brainf\*ck Parsing Rules

file: instructions EOF;

instructions: instruction+;

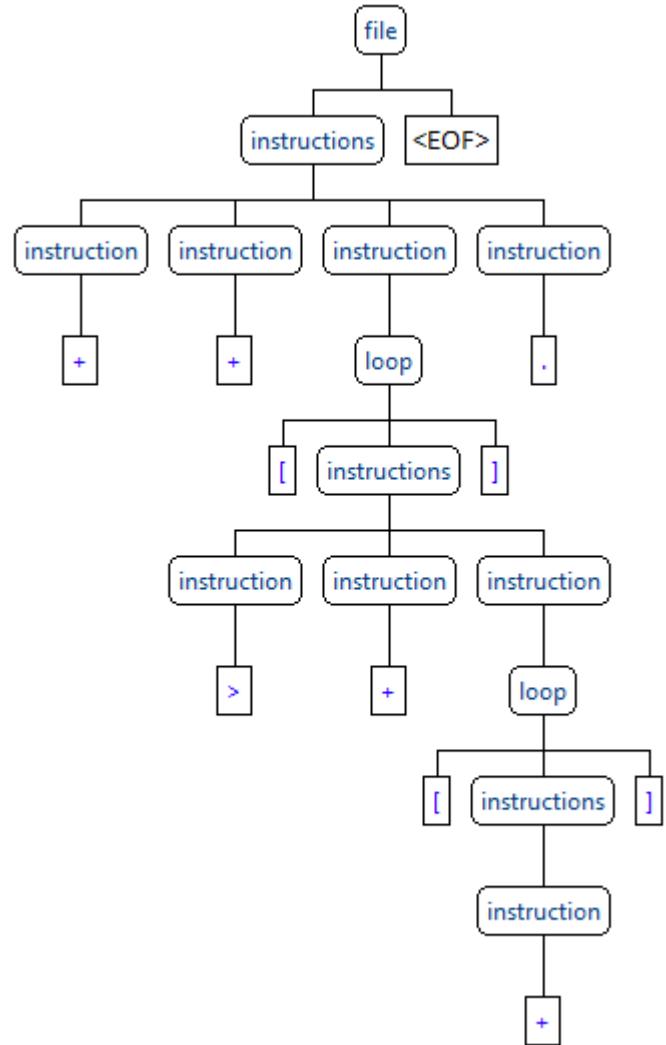
instruction: loop  
| MOVE\_RIGHT  
| MOVE\_LEFT  
| INCREMENT  
| DECREMENT  
| WRITE  
| READ  
;

loop: LOOP\_HEADER instructions+ LOOP\_FOOTER;

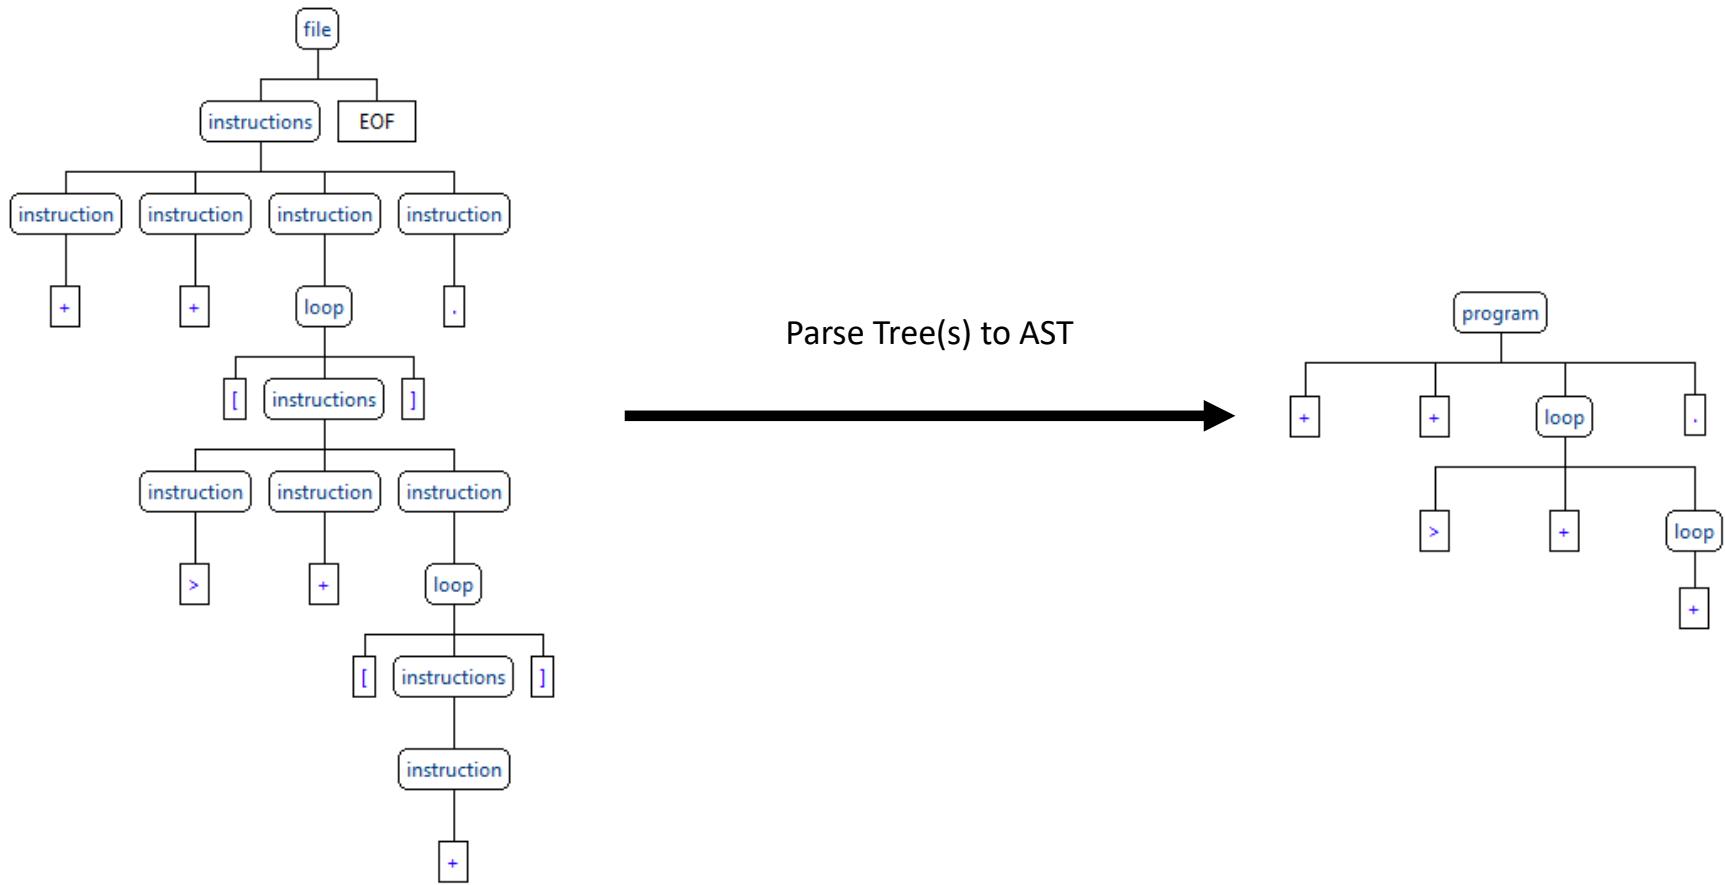


# Brainf\*ck Parse Tree

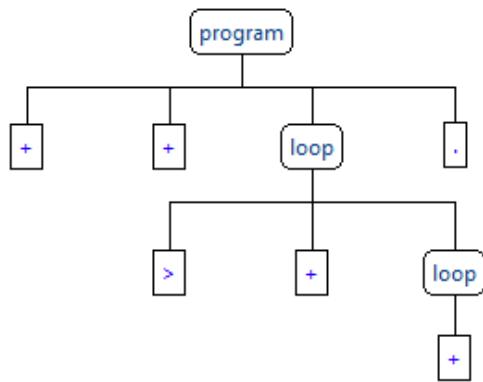
Program: `++>+[+]`.



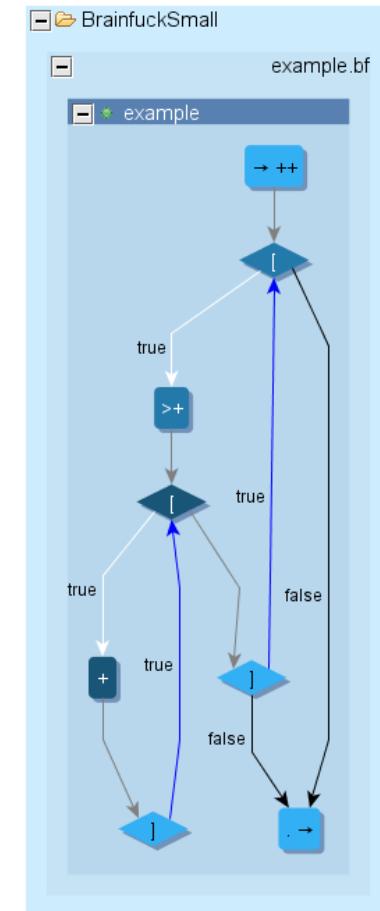
# Brainf\*ck Abstract Syntax Tree (AST)



# Brainf\*ck AST to Program Graph



Parse Tree(s) to AST



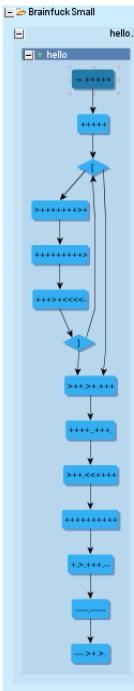
• Brainf\*ck Hello World Program

• Graph contains information necessary to execute program

• This language should be simple to analyze right???

- No variables, just tape cells
- How many behaviors could there be?

```
hello.bf
1 +++++
2 +++++
3 [>++++++>++++++>+++><--]
4 >++.>+.+++++.++.>+.<+++++++.>.++.-----.>+,.
```



Operation completed successfully...

```
show(cfg(universe.nodes(XCSG.Function)))
com.benjholla.atlas.brainfuck.interpreter.BrainfuckGraphInterpreter.execute(selected)
res1: String =
"Hello World!
"
Evaluate: |
```

# Elemental: A Brainf\*ck Derivative

- [github.com/benjholla/Elemental](https://github.com/benjholla/Elemental)
  - Goal is to be basic, not to be tiny
  - Separates looping and branching
  - New features to explore impacts of modern language features

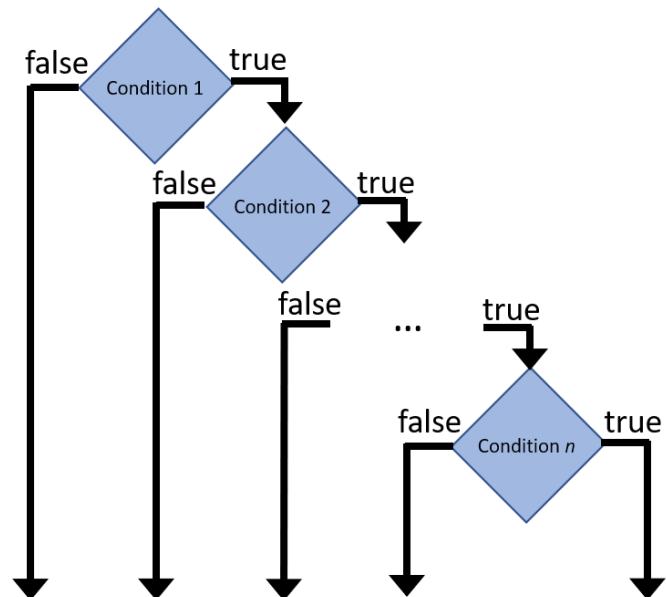
Instruction	Description
+	Increment the byte at the current tape cell by 1
-	Decrement the byte at the current tape cell by 1
<	Move the tape one cell to the left
>	Move the tape one cell to the right
,	(Store) Read byte value from input into current tape cell
.	(Recall) Write byte value to output from current tape cell
(	(Branch) If the byte value at the current cell is 0 then jump to the instruction following the matching ), else execute the next instruction
[	(While Loop) If the byte value at the current cell is 0 then jump to the instruction following the matching ], else execute instructions until the matching ] and then unconditionally return to the [
[0-9]+:	(Function) Declares a uniquely named function (named [0-9]+ within range 0-255)
{[0-9]+}	(Static Dispatch) Jump to a named function
?	(Dynamic Dispatch/Function Pointer) Jumps to a named function with the value of the current cell
"[0-9]+"	(Label) Sets a unique label (named [0-9]+ within range 0-255) within a function
'[0-9]+'	(GOTO) Jumps to a named label within the current function
&	(Computed GOTO) Jumps to the named label within the current function with the value of the current cell
#	A one line comment

# Counting Program Paths: Branching

How many paths are there for  $n$  nested branches?

```
if(condition_1){  
    if(condition_2){  
        if(condition_3){  
            ...  
            if(condition_n){  
                // conditions 1 through n  
                // must all be true to reach here  
            }  
        }  
    }  
}
```

**$n+1$  paths!**

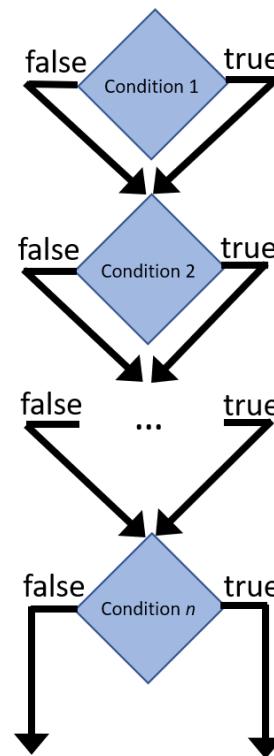


# Counting Program Paths: Branching

How many paths are there for  $n$  non-nested branches?

```
if(condition_1){  
    // code block 1  
}  
if(condition_2){  
    // code block 2  
}  
if(condition_3){  
    // code block 3  
}  
...  
if(condition_n){  
    // code block n  
}
```

$2^n$  paths!



# Elemental: A Brainf\*ck Derivative

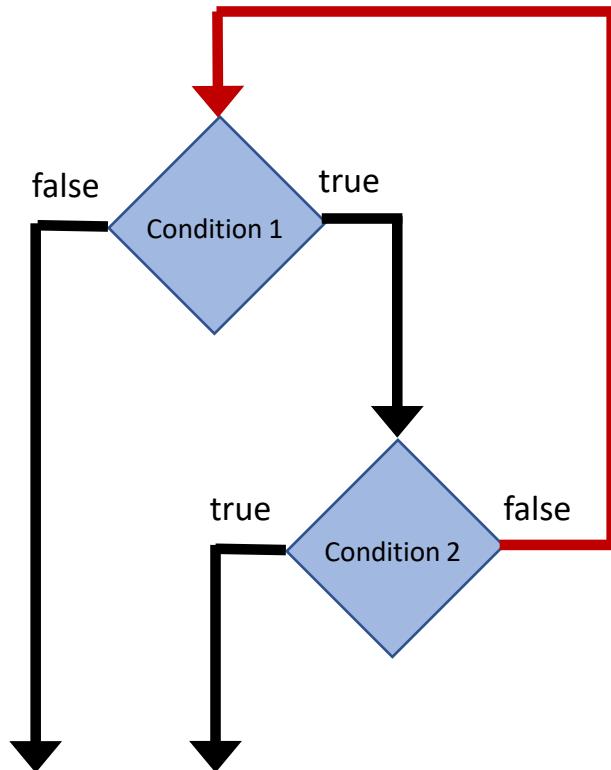
- [github.com/benjholla/Elemental](https://github.com/benjholla/Elemental)
  - Goal is to be basic, not to be tiny
  - Separates looping and branching
  - New features to explore impacts of modern language features

Instruction	Description
+	Increment the byte at the current tape cell by 1
-	Decrement the byte at the current tape cell by 1
<	Move the tape one cell to the left
>	Move the tape one cell to the right
,	(Store) Read byte value from input into current tape cell
.	(Recall) Write byte value to output from current tape cell
(	(Branch) If the byte value at the current cell is 0 then jump to the instruction following the matching ), else execute the next instruction
[	(While Loop) If the byte value at the current cell is 0 then jump to the instruction following the matching ], else execute instructions until the matching ] and then unconditionally return to the [
[0-9]+:	(Function) Declares a uniquely named function (named [0-9]+ within range 0-255)
{[0-9]+}	(Static Dispatch) Jump to a named function
?	(Dynamic Dispatch/Function Pointer) Jumps to a named function with the value of the current cell
"[0-9]+"	(Label) Sets a unique label (named [0-9]+ within range 0-255) within a function
'[0-9]+'	(GOTO) Jumps to a named label within the current function
&	(Computed GOTO) Jumps to the named label within the current function with the value of the current cell
#	A one line comment

# Considering Loops

- Programs may have loops
  - How many paths does this program have?
  - Can we say if this program halts?

```
while(condition_1){  
    if(condition_2){  
        break;  
    }  
}
```



# The Halting Problem

Suppose, we could construct:

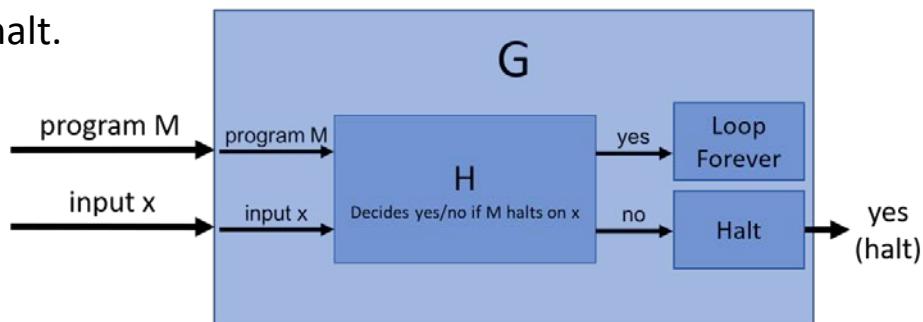
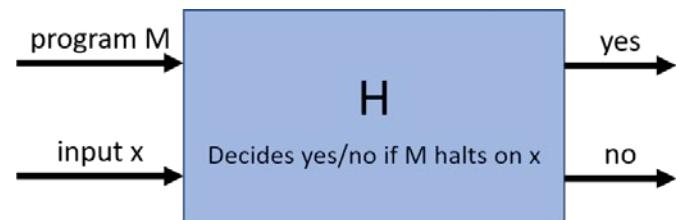
$H(M, x) :=$  if  $M$  halts on  $x$  then return true else return false

Then we could construct:

$G(M, x) :=$  if  $G(M, x)$  is false then return true else loop forever

But if we then pass  $G$  to itself, that is  $G(G, G)$ , we get a contradiction between what  $G$  does and what  $H$  says that  $G$  does. If  $H$  says that  $G$  halts, then  $G$  does not halt. If  $H$  says that  $G$  does not halt, then it does halt.

$H$  cannot exist.



# Elemental: A Brainf\*ck Derivative

- [github.com/benjholla/Elemental](https://github.com/benjholla/Elemental)
  - Goal is to be basic, not to be tiny
  - Separates looping and branching
  - New features to explore impacts of modern language features
- '?' could pass control to any function!
- '&' could jump to any line!
- Goto labels with '?' or '&' could be simulated with branching or loops
- These blur control flow with data

Instruction	Description
+	Increment the byte at the current tape cell by 1
-	Decrement the byte at the current tape cell by 1
<	Move the tape one cell to the left
>	Move the tape one cell to the right
,	(Store) Read byte value from input into current tape cell
.	(Recall) Write byte value to output from current tape cell
(	(Branch) If the byte value at the current cell is 0 then jump to the instruction following the matching ), else execute the next instruction
[	(While Loop) If the byte value at the current cell is 0 then jump to the instruction following the matching ], else execute instructions until the matching ] and then unconditionally return to the [
[0-9]+:	(Function) Declares a uniquely named function (named [0-9]+ within range 0-255)
{[0-9]+}	(Static Dispatch) Jump to a named function
?	(Dynamic Dispatch/Function Pointer) Jumps to a named function with the value of the current cell
"[0-9]+"	(Label) Sets a unique label (named [0-9]+ within range 0-255) within a function
'[0-9]+'	(GOTO) Jumps to a named label within the current function
&	(Computed GOTO) Jumps to the named label within the current function with the value of the current cell
#	A one line comment

# Positive Trend – Address the Languages

- Data drives execution
  - Data is half of the program!
  - “The illusion that your program is manipulating its data is powerful. But it is an illusion: The data is controlling your program.”
- Crema: A LangSec-Inspired Programming Language
  - Giving a developer a Turning complete language for every task is like *giving a 16 year old a formula one car* (something bad is bound to happen soon)



# milcom

Military Communications for the 21st Century  
October 29-31, 2018 • LAX Marriott, Los Angeles, CA

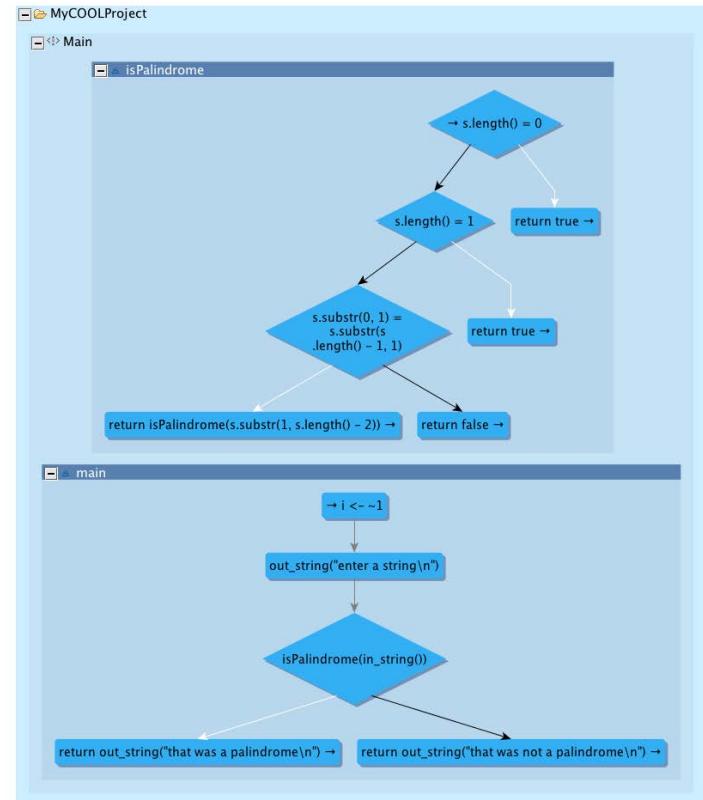


# Positive Trend – Address the Languages

- Data drives execution
  - Data is half of the program!
  - “The illusion that your program is manipulating its data is powerful. But it is an illusion: The data is controlling your program.”
- Crema: A LangSec-Inspired Programming Language
  - Giving a developer a Turning complete language for every task is like *giving a 16 year old a formula one car* (something bad is bound to happen soon)
  - Apply principle of least privilege to computation (least computation principle)
    - Computational power exposed to attacker *is* privilege. Minimize it.
    - Try copy-pasting the XML billion-laughs attack from Notepad into MS Word if you want to see why...

# Scaling Up: Program Analysis for COOL

- Classroom Object Oriented Language (COOL)
  - [https://en.wikipedia.org/wiki/Cool\\_\(programming\\_language\)](https://en.wikipedia.org/wiki/Cool_(programming_language))
  - <http://openclassroom.stanford.edu/MainFolder/CoursePage.php?course=Compilers>
- COOL Program Graph Indexer
  - Type hierarchy
  - Containment relationships
  - Function / Global variable signatures
  - Function Control Flow Graph
  - Data Flow Graph (in progress)
  - Inter-procedural relationships:
    - Call Graph (implemented via compliance to XCSG!)
- <https://github.com/benjholla/AtlasCOOL>  
(currently private)



# Program Analysis for Contemporary Languages

- <http://www.ensoftcorp.com/atlas> (Atlas)
  - C, C++, Java Source, Java Bytecode, *and now Brainfuck/COOL!*
- <https://scitools.com> (Understand)
  - C, C++ Source
- <http://mlsec.org/joern> (Joern)
  - C, C++, PHP Source
- <https://www.hex-rays.com/products/ida> (IDA)
- <https://binary.ninja> (Binary Ninja)
- <https://www.radare.org> (Radare)

# Data Flow Graph (DFG)

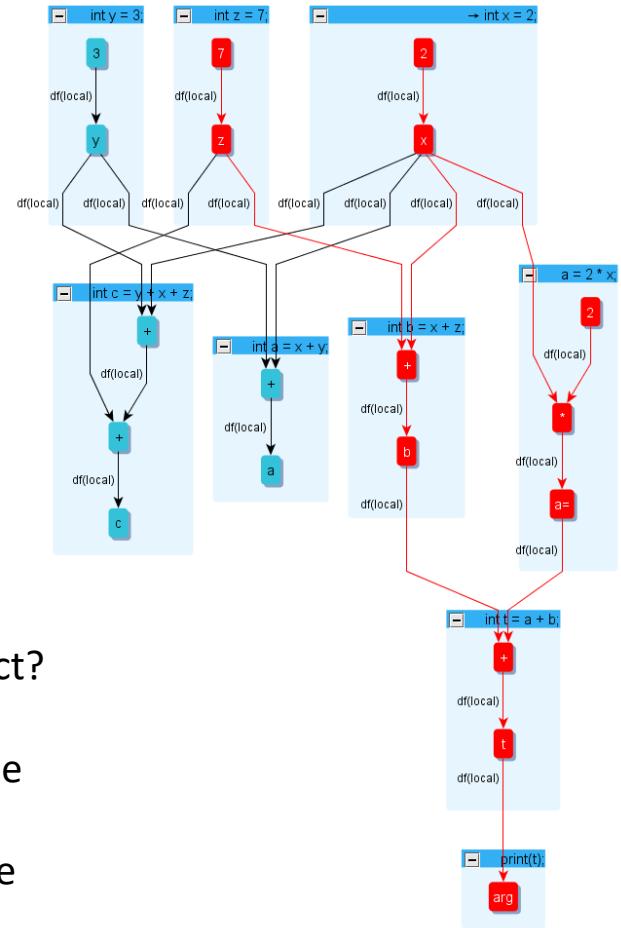
Example:

1.  $x = 2;$
2.  $y = 3;$
3.  $z = 7;$
4.  $a = x + y;$
5.  $b = x + z;$
6.  $a = 2 * x;$
7.  $c = y + x + z;$
8.  $t = a + b;$
9.  $\text{print}(t);$

Relevant lines:  
1,3,5,6,8

detected failure

- What lines must we consider if the value of  $t$  printed is incorrect?
- A *Data Flow Graph* creates a graph of primitives and variables where each assignment represents an edge from the RHS to the LHS of the assignment
- The *Data Flow Graph* represents global data dependence at the operator level (the atomic level) [FOW87]



# Control Flow Graph (CFG)

Example:

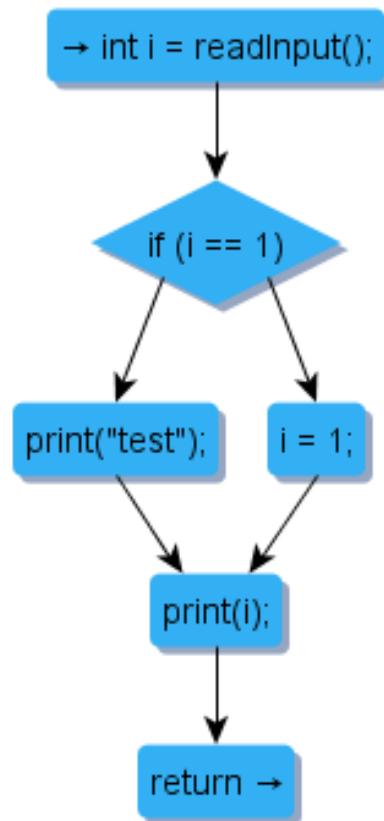
```

1. i = readInput();
2. if(i == 1)
3.     print("test");
    else
4.     i = 1;
5. print(i); ← detected failure
6. return; // terminate
  
```

Relevant lines:  
1,2,4

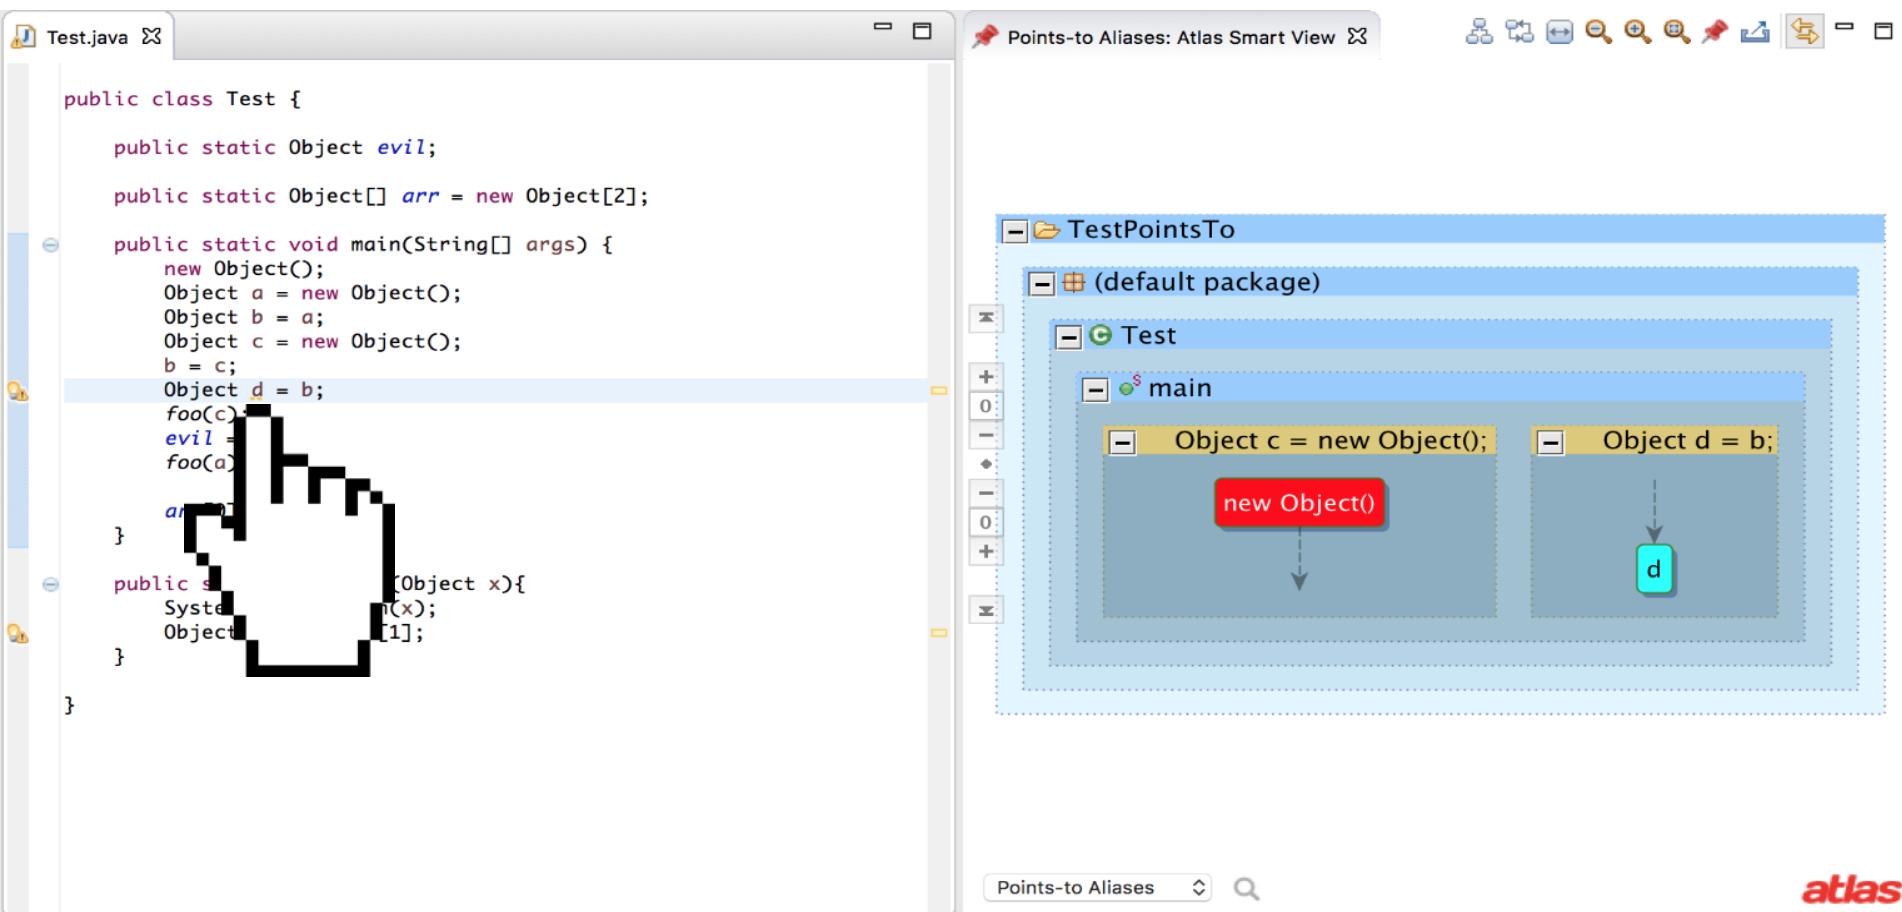
What lines must we consider if the value of *i* printed is incorrect?

- A *Control Flow Graph (CFG)* represents the possible sequential execution orderings of each statement in a program
- Data flow influences control flow, so this graph is not enough



# Analysis Woes: Going Inter-procedural

- Function pointers and dynamic dispatches force us to solve data flow problems to precisely identify inter-procedural control flows
- Class Hierarchy Analysis / Rapid Type Analysis
  - Cheap but *very* conservative results
- Points-to Analysis
  - A variable  $v$  points-to what data in memory? Knowing this we can more precisely resolve
  - Obtaining a perfect solution has been proven to reduce to solving the halting problem...
  - Expensive even for conservative results!!!
    - Each level of precision adds an exponent
    - Not really a lot to be gained (in our opinion)

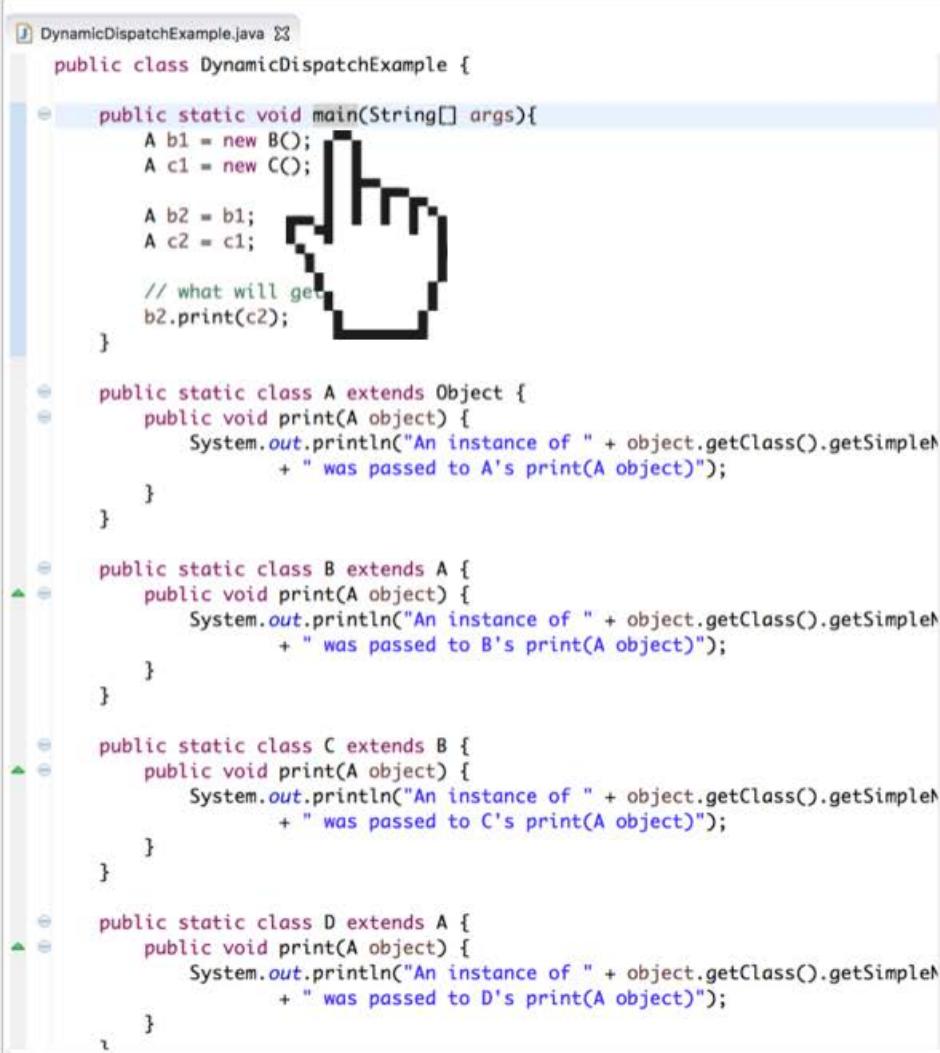


The screenshot shows a Java development environment with two main windows:

- Test.java**: A code editor window displaying the following Java code:

```
public class Test {  
    public static Object evil;  
    public static Object[] arr = new Object[2];  
  
    public static void main(String[] args) {  
        new Object();  
        Object a = new Object();  
        Object b = a;  
        Object c = new Object();  
        b = c;  
        Object d = b;  
        foo(c);  
        evil =  
        foo(a);  
  
        a =  
    }  
  
    public static void foo(Object x){  
        System.out.println(x);  
    }  
}
```

A large black mouse cursor is positioned over the code editor.
- Points-to Aliases: Atlas Smart View**: An analysis tool window titled "TestPointsTo". It shows a call graph for the "main" method of the "Test" class. The graph highlights several nodes:
  - A red box highlights the statement `Object c = new Object();`.
  - A yellow box highlights the statement `Object d = b;`.
  - A blue box highlights the variable `d`.
  - A red box highlights the expression `new Object()`.The analysis tool interface includes a sidebar with icons for file operations, a search bar at the bottom, and the "atlas" logo in the bottom right corner.



```

DynamicDispatchExample.java

public class DynamicDispatchExample {

    public static void main(String[] args){
        A b1 = new B();
        A c1 = new C();

        A b2 = b1;
        A c2 = c1;

        // what will get
        b2.print(c2);
    }

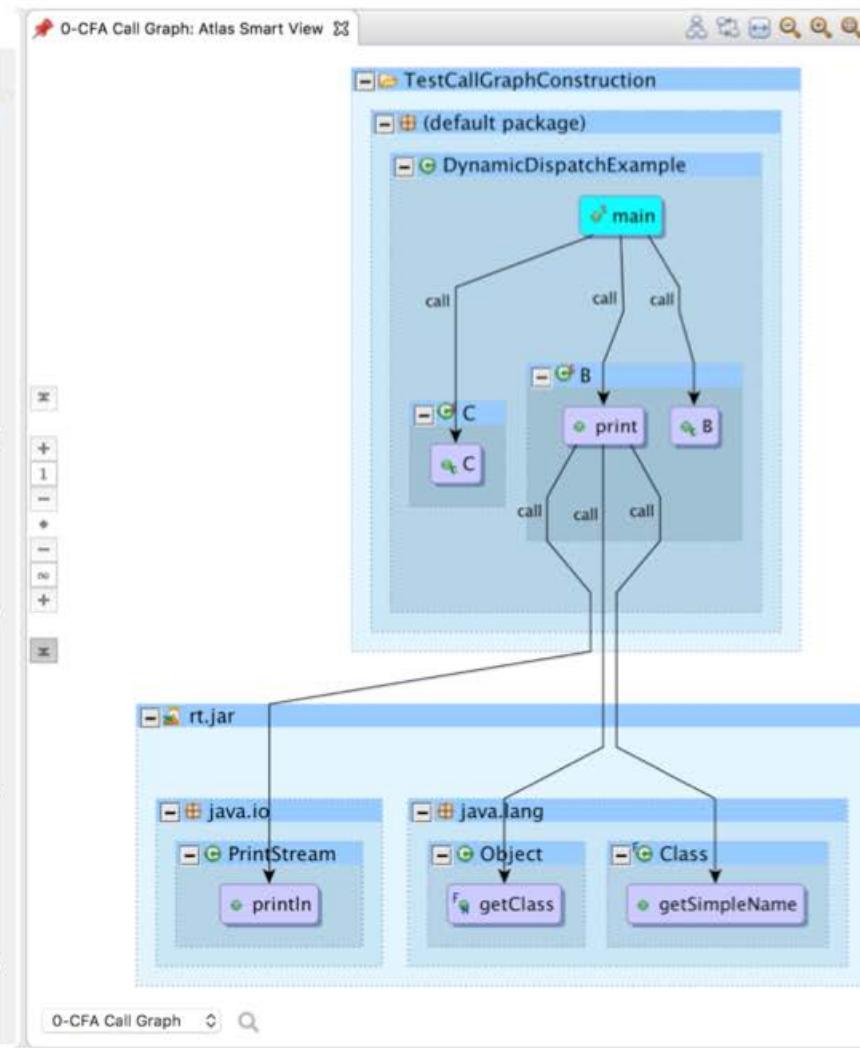
    public static class A extends Object {
        public void print(A object) {
            System.out.println("An instance of " + object.getClass().getSimpleName()
                + " was passed to A's print(A object)");
        }
    }

    public static class B extends A {
        public void print(A object) {
            System.out.println("An instance of " + object.getClass().getSimpleName()
                + " was passed to B's print(A object)");
        }
    }

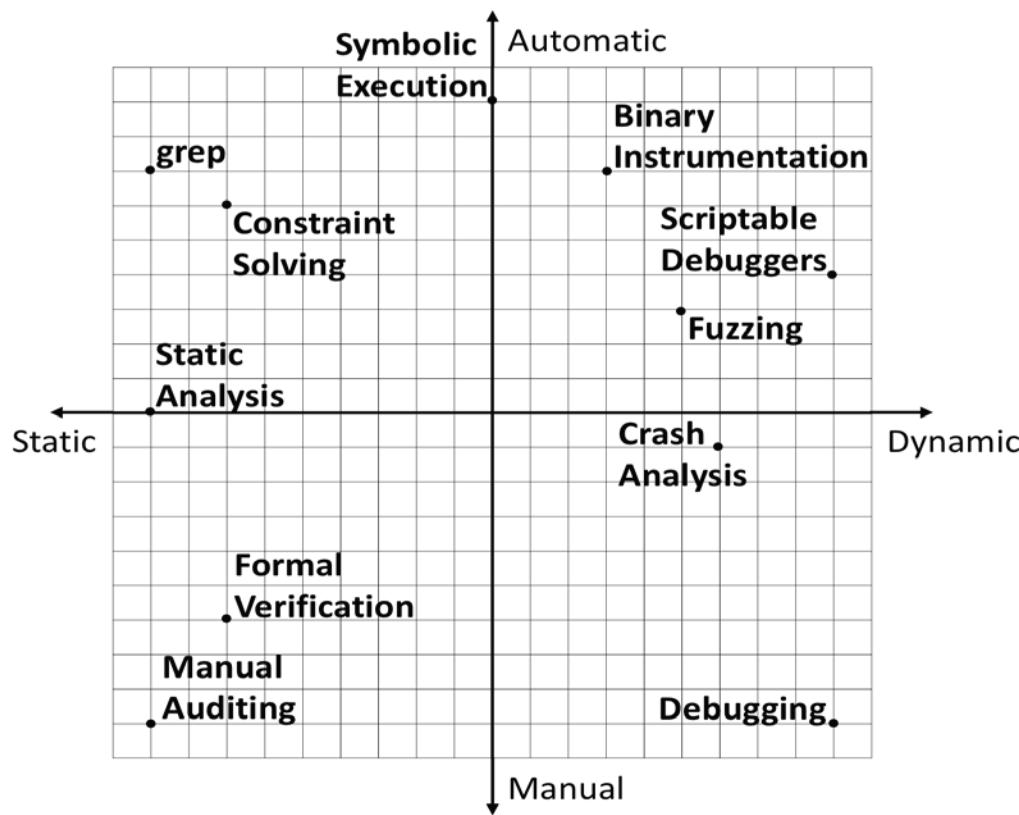
    public static class C extends B {
        public void print(A object) {
            System.out.println("An instance of " + object.getClass().getSimpleName()
                + " was passed to C's print(A object)");
        }
    }

    public static class D extends A {
        public void print(A object) {
            System.out.println("An instance of " + object.getClass().getSimpleName()
                + " was passed to D's print(A object)");
        }
    }
}

```



# A Spectrum of Program Analysis Techniques



Source: Contemporary Automatic Program Analysis,  
Julian Cohen, Blackhat 2014

# Symbolic Execution

- Replace concrete assignment values with symbolic values
- Perform operations on symbolic values abstractly
- At each branch fork the abstracted logic
  - Dealing with path explosion problem is a challenge
- Utilize SAT/SMT solvers to determine if the constraints are satisfiable for a path of interest
  - Example: fail occurs if  $y * 2 = z = 12$  is satisfiable
    - Solve( $y * 2 = 12$ ,  $y$ ),  $y = 6$  satisfies the constraint
    - Program crash occurs when `read()` returns 6

```
int f() {  
    ...  
    y = read();  
    z = y * 2;  
    if (z == 12) {  
        crash();  
    } else {  
        printf("OK");  
    }  
}
```

On what inputs does the code crash?

```
#include <stdio.h>
#include <stdlib.h>

char *serial = "\x31\x3e\x3d\x26\x31";

int check(char *ptr)
{
    int i;
    int hash = 0xABCD;

    for (i = 0; ptr[i]; i++)
        hash += ptr[i] ^ serial[i % 5];

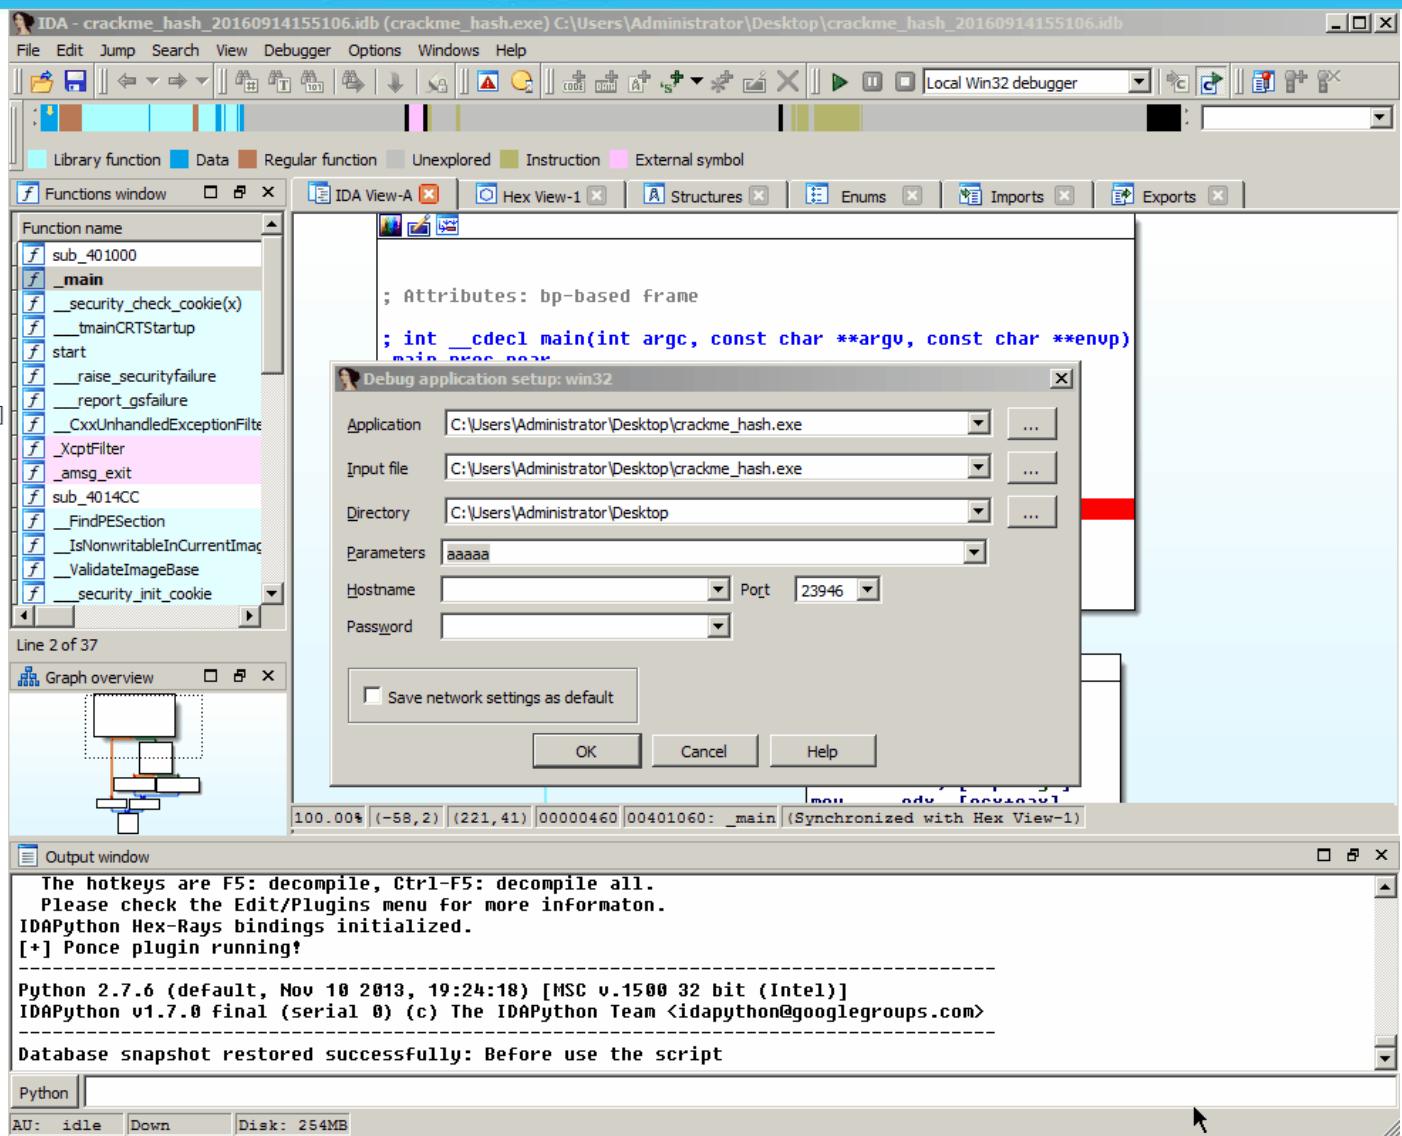
    return hash;
}

int main(int ac, char **av)
{
    int ret;

    if (ac != 2)
        return -1;

    ret = check(av[1]);
    if (ret == 0xad6d)
        printf("Win\n");
    else
        printf("fail\n");

    return 0;
}
```

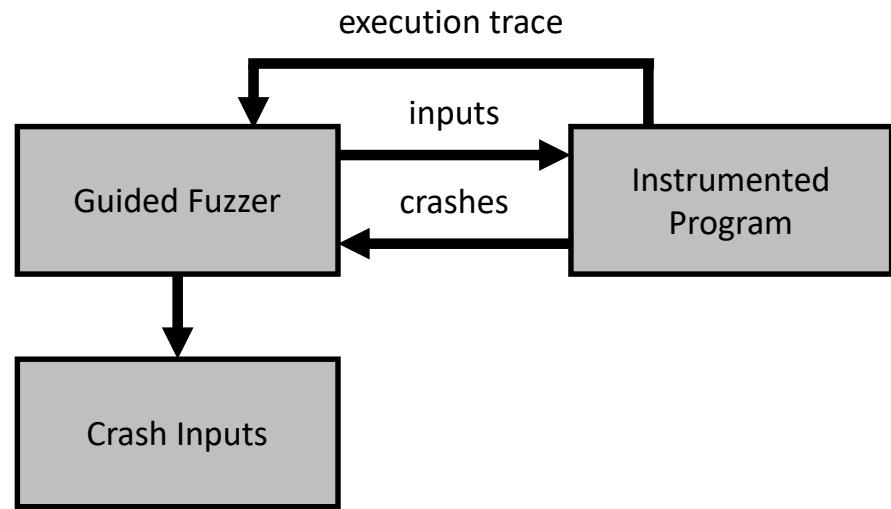


# Concolic Execution

- A hybrid of dynamic analysis and symbolic execution
  - Perform symbolic execution on some variables and concrete execution on other variables
  - Symbolic variables could be made concrete in order to:
    - Move past symbolic limitations such as challenges in modeling the program environment (example network interaction)
    - Deal with path explosion problem and satisfiability problem by replacing difficult symbolic values with concrete values to simplify analysis
  - Pays cost in time for symbolic computations and execution time of program
- Several well known tools:
  - Angr - <http://angr.io>
  - KLEE - <https://klee.github.io>
  - DART - <https://dl.acm.org/citation.cfm?id=1065036>
  - CREST (formerly CUTE) - <https://code.google.com/archive/p/crest>
  - Microsoft SAGE - [https://patricegodefroid.github.io/public\\_psfiles/ndss2008.pdf](https://patricegodefroid.github.io/public_psfiles/ndss2008.pdf)

# Smart (Guided) Fuzzing

- Start with a test corpus of well formed program inputs
- Apply random or systematic mutations to program inputs
- Instrument the program branch points
- Run the instrumented program with mutated inputs and 1) observe whether or not the program crashes and 2) record the program execution path coverage
- If the input results in new program paths being explored then prioritize mutations of the tested input
- Repeat until the program crashes



Heuristics guide genetic algorithm to generate program inputs that push the fuzzer deeper into the program control flow, avoiding the common pitfalls of fuzzers to only test “shallow” code regions.

# AFL (American Fuzzy Lop) Fuzzer

- Recognized as the current state of art implementation of guided fuzzing
  - Effective mutation strategy to generate new inputs from initial test corpus
  - Lightweight instrumentation at branch points
  - Genetic algorithm promotes mutations of inputs that discover new branch edges
    - Aims to explore all code paths
  - Huge trophy case of bugs found in wild
    - 371+ reported bugs in 161 different programs as of March 2018
    - <http://lcamtuf.coredump.cx/afl/>
- A game of economics. AFL tends to “guess” the correct input faster than a smart tool “computes” the correct input.

```
american fuzzy lop 0.47b (readpng)

process timing    run time : 0 days, 0 hrs, 4 min, 43 sec          overall results
last new path   : 0 days, 0 hrs, 0 min, 26 sec      cycles done : 0
last uniq crash  : none seen yet                  total paths : 195
last uniq hang   : 0 days, 0 hrs, 1 min, 51 sec     uniq crashes : 0
                                         map coverage
cycle progress  now processing : 38 (19.49%)      map density : 1217 (7.43%)
                paths timed out : 0 (0.00%)       count coverage : 2.55 bits/tuple
stage progress  now trying : interest 32/8        findings in depth
                stage execs : 0/9990 (0.00%)      favored paths : 128 (65.64%)
                                         new edges on : 85 (43.59%)
                                         total execs : 654k
                                         exec speed : 2306/sec      total crashes : 0 (0 unique)
                                         total hangs : 1 (1 unique)
                                         path geometry
fuzzing strategy yields
bit flips        : 88/14.4k, 6/14.4k, 6/14.4k
byte flips       : 0/1804, 0/1786, 1/1750
arithmetics     : 31/126k, 3/45.6k, 1/17.8k
known ints       : 1/15.8k, 4/65.8k, 6/78.2k
havoc           : 34/254k, 0/0
trim             : 2876 B/931 (61.45% gain)
                                         levels : 3
                                         pending : 178
                                         pend fav : 114
                                         imported : 0
                                         variable : 0
                                         latent : 0
```

# DARPA's Cyber Grand Challenge (CGC)

“Cyber Grand Challenge (CGC) is a contest to build high-performance computers capable of playing in a Capture-the-Flag style cyber-security competition.”



# DARPA's Cyber Grand Challenge (CGC)

- Fully automatic reasoning to:
  - Detect program vulnerabilities
  - Patch programs to prevent exploitation
  - Develop and execute vulnerability exploits against competitors
- No human players!



# CGC Results (Reading Between the Lines)

- All teams published the same essential combination of strategies
  - Guided fuzzing (nearly every team modified AFL)
  - Symbolic execution to assist fuzzer sometimes aided by classical program analyses (points-to, reachability, slicing, etc.)
  - Some state space pruning and prioritization scheme catered to expected vulnerability types
- Effective patches were more often generic patches which addressed the class of vulnerabilities not the one-off vulnerability that was given
  - Example: Adding stack guards for memory protection
- Competitor scores were close!
  - The difference between 1<sup>st</sup> and 7<sup>th</sup> place was not substantial
- Classes of vulnerabilities were known *a priori*

# Context Matters!

- Head problems vs. hand problems
  - Design vs. implementation errors
- Implementation errors may be eventually largely eliminated with advances in programming languages, compilers, theorem provers, etc.
- Security design problems tend to be closely tied to failures in threat modeling, which is largely a human task
  - Many security problems arise due to reuse of software in changing contexts
    - Example: SCADA devices designed for isolated networks being placed on the internet

# DARPA's High-Assurance Cyber Military Systems (HACMS) Program

- “formal methods-based approach to enable semi-automated code synthesis from executable, formal specifications”
- Creation of an “unhackable” drone!

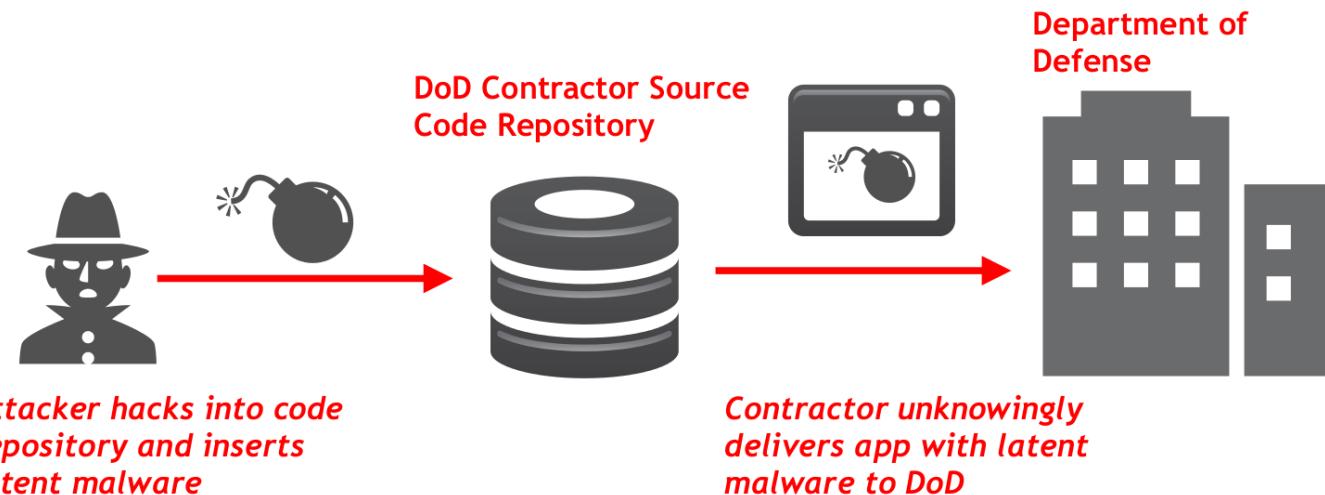


# HACMS Results (Reading Between the Lines)

- Failures tended to stem back to human failures to fully account for the attack model
  - Example: Red Team debrief noted that Red team sent a radio reboot command that dropped that shut off drone engines for 3 seconds (enough to crash the drone). Blue's response was "Oh! We didn't think of that!"
- The "unhackable" drone produced by the program was not protected from the later discovered Meltdown and Spectre exploits

# DARPA's APAC Program

- Automated Program Analysis For Cybersecurity (APAC)
- Scenario: Hardened devices, internal app store, untrusted contractors, expert adversaries
- Focused on Android



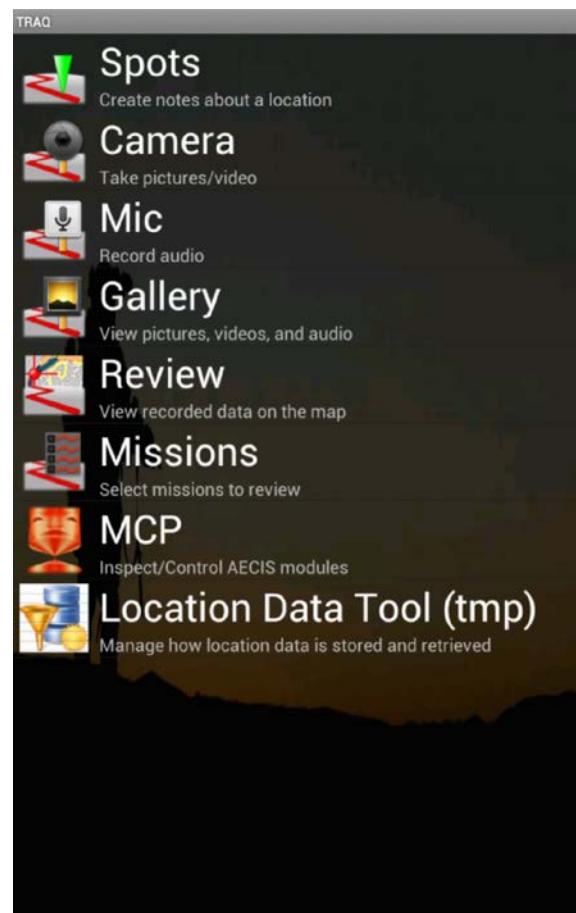
# APAC Results (Reading Between the Lines)

- Is a bug malware? What if its planted intentionally? We know bug finding is hard...
- Bugs have plausible deniability and malicious intent cannot be determined from code.
- Novel attacks have escaped previous threat models.
- Need precision tools to detect **novel** and **sophisticated** malware in advance!



# APAC Example: DARPA's Transformative Apps

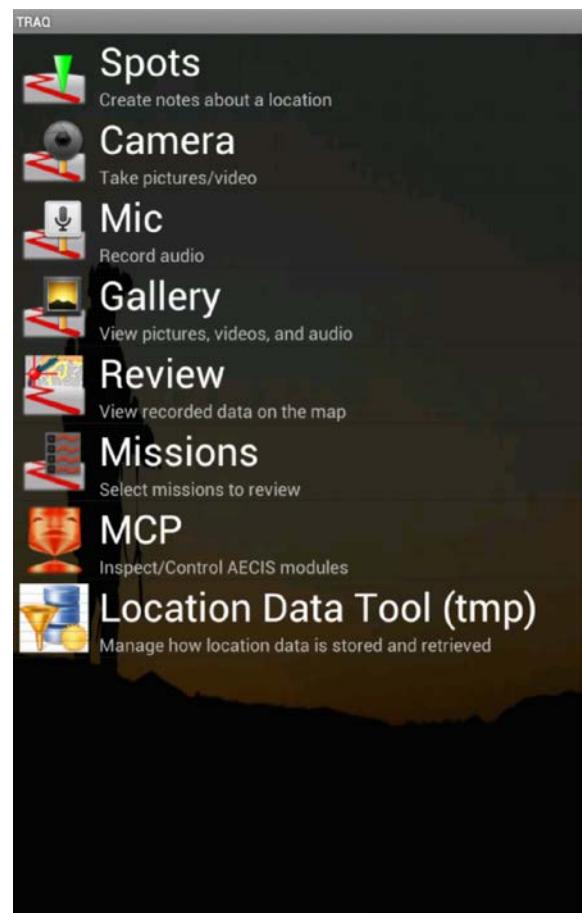
- 55K lines of code
- Data gathering and relaying tool for military
  - Strategic mission planning/review
  - Audio and video recording
  - Geo-tagged camera snapshots
  - Real-time map updates based on GPS



# APAC Example: DARPA's Transformative Apps

```
@Override
public void onLocationChanged(Location tmpLoc) {
    location = tmpLoc;
    double latitude = location.getLatitude();
    double longitude = location.getLongitude();
    if((longitude >= 62.45 && longitude <= 73.10) &&
       (latitude >= 25.14 && latitude <= 37.88)) {
        location.setLongitude(location.getLongitude() + 9.252);
        location.setLatitude(location.getLatitude() + 5.173);
    }
    ...
}
```

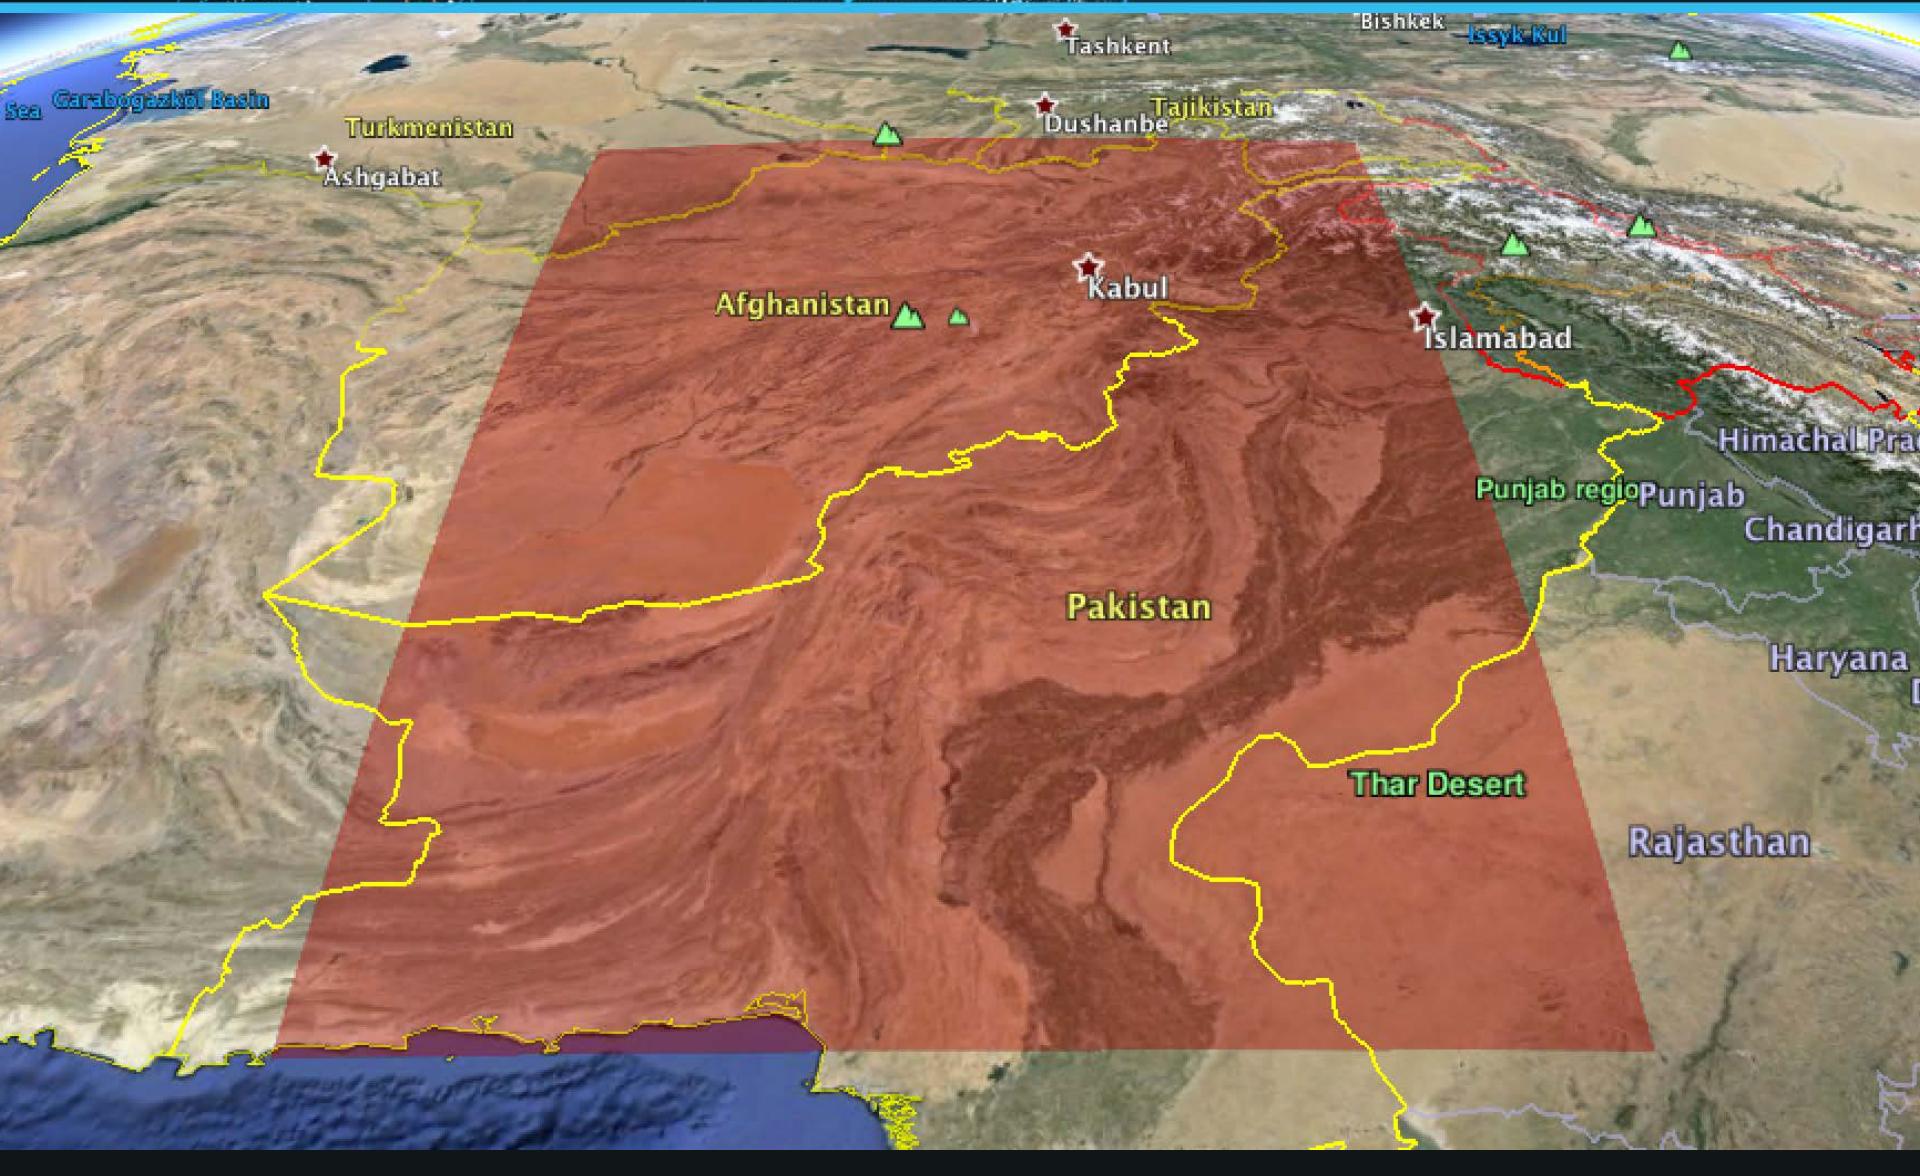
One of several locations were GPS coordinates were modified.  
This corrupts GPS coordinates if user is in Afghanistan/Pakistan!





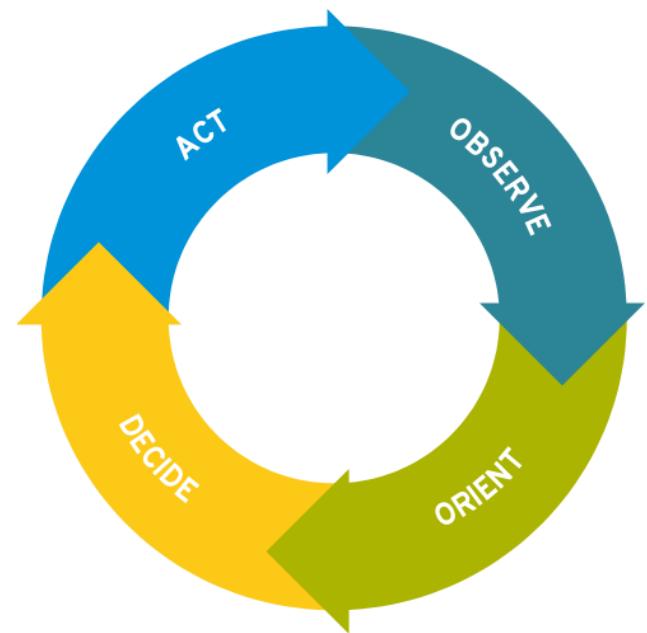
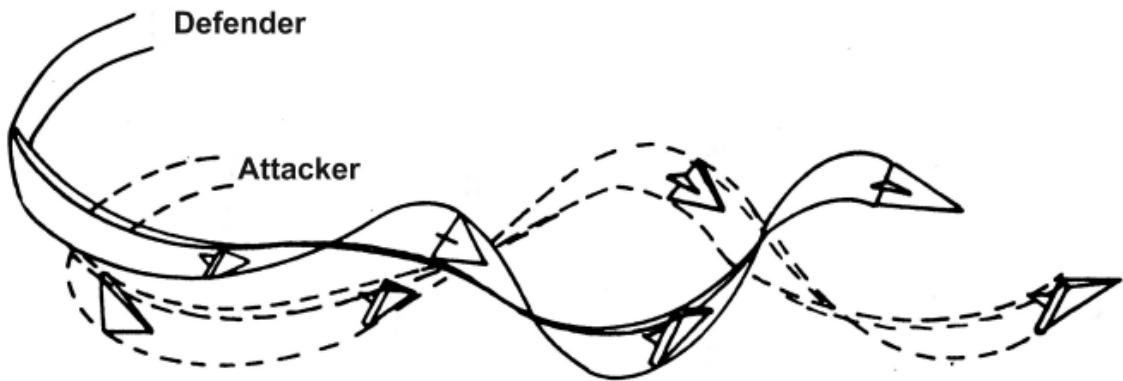
# milcom

Military Communications for the 21st Century  
October 29-31, 2018 • LAX Marriott, Los Angeles, CA

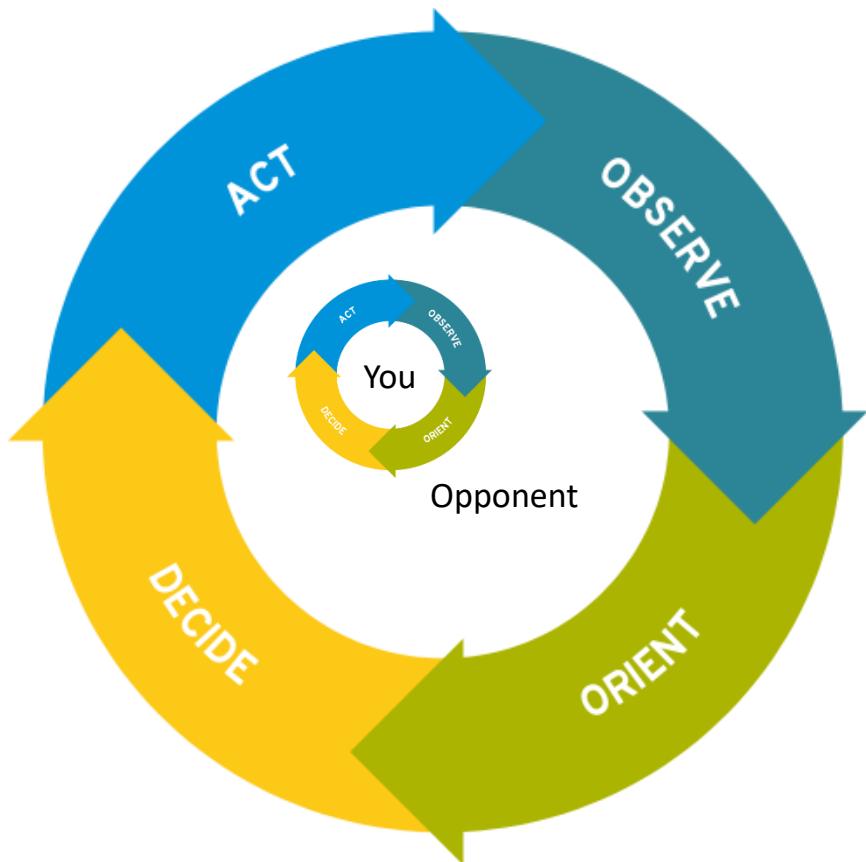


# Program Analysis, OODA, and YOU

“Security is a process, not a product” – Bruce Schneier



# Program Analysis, OODA, and YOU



Our opponent

- Time
- Evolution of malware

“...IA > AI, that is, that intelligence amplifying systems can, at any given level of available systems technology, beat AI systems. That is, a machine and a mind can beat a mind-imitating machine working by itself.”  
– Fred Brooks

# Positive Trend – Employ Human Reasoning

- DARPA Programs on the verge
  - Computers and Humans Exploring Software Security (CHESS)
  - Explainable Artificial Intelligence (XAI)
- Let human's do what humans are good at
- Let machines do what machines are good at
- Enable collaboration between the two



Military Communications for the 21st Century  
October 29-31, 2018 • LAX Marriott, Los Angeles, CA

Your job in security is not going away soon...

Your job in security is not going away soon...  
...probably.

(but it is an exciting field)

## Activity: Does this program contain a vulnerability?

```
#include <stdio.h>
int main(int argc, char *argv) {
    char buf[64];
    strcpy(buf, argv[1]);
    return 0;
}
```

## Activity: Does this program contain a vulnerability?

```
#include <stdio.h>
int main(int argc, char *argv) {
    char buf[64];
    a = read_user_input();
    b = read_user_input();
    ...
    if(a * b == c){
        strcpy(buf, argv[1]);
    }
    return 0;
}
```

What if  $c$  is the product of two large primes?

Then yes, but only if you know the prime factorization...

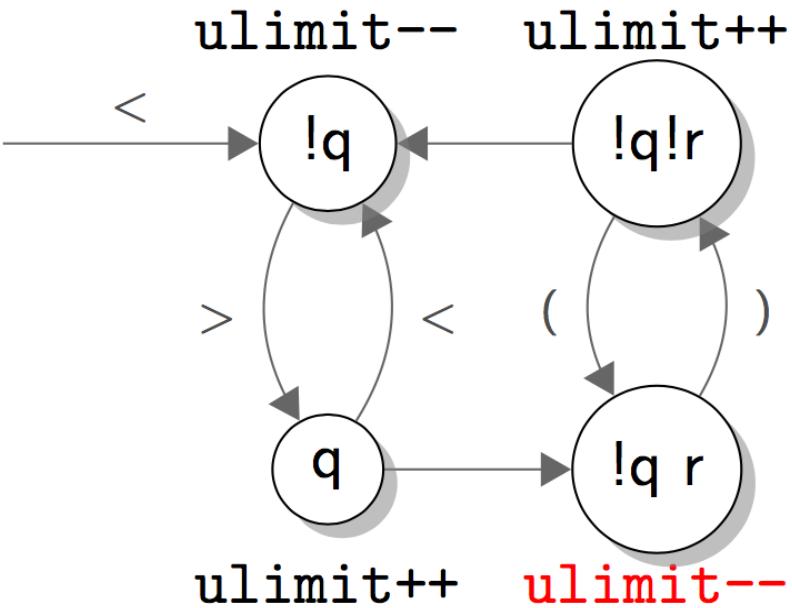
# Activity: Does this program contain a vulnerability?

```
#define BUFFERSIZE 200
int copy_it (char* input , unsigned int length){
    char c, localbuf[BUFFERSIZE];
    unsigned int upperlimit = BUFFERSIZE - 10;
    unsigned int quotation = roundquote = FALSE;
    unsigned int input_index = output_index = 0;
    while (input_index < length){ c = input[input_index++];
        if((c == '<') && (!quotation)){ quotation = true; upperlimit--; }
        if((c == '>') && (quotation)){ quotation = false; upperlimit++; }
        if((c == '(') && (!quotation) && (!roundquote)){ roundquote = true; }
        if((c == ')') && (!quotation) && (roundquote)){ roundquote = false; upperlimit++; }
        //if there is sufficient space in the buffer, write the character
        if(output_index < upperlimit){ localbuf[output_index] = c; output_index++; }
    }
    if(roundquote){ localbuf[output_index] = ')'; output_index++; }
    if(quotation){ localbuf[output_index] = '>'; output_index++; }
    return output_index;
}
```

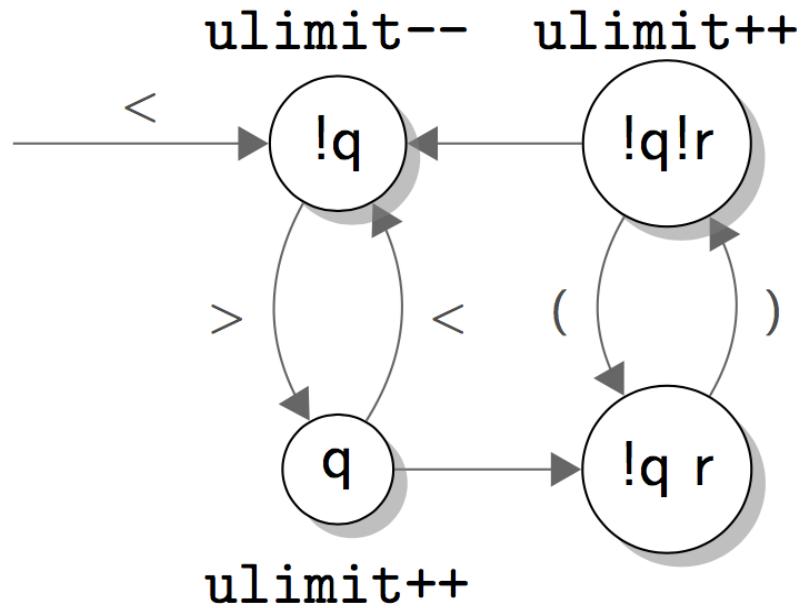
# Activity: Does this program contain a vulnerability?

```
#define BUFFERSIZE 200
int copy_it (char* input , unsigned int length){
    char c, localbuf[BUFFERSIZE];
    unsigned int upperlimit = BUFFERSIZE - 10;
    unsigned int quotation = roundquote = FALSE;
    unsigned int input_index = output_index = 0;
    while (input_index < length){ c = input[input_index++];
        if((c == '<') && (!quotation)){ quotation = true; upperlimit--; }
        if((c == '>') && (quotation)){ quotation = false; upperlimit++; }
        if((c == '(') && (!quotation) && (!roundquote)){ roundquote = true; /* (missing) upperlimit--; */ }
        if((c == ')') && (!quotation) && (roundquote)){ roundquote = false; upperlimit++; }
        //if there is sufficient space in the buffer, write the character
        if(output_index < upperlimit){ localbuf[output_index] = c; output_index++; }
    }
    if(roundquote){ localbuf[output_index] = ')'; output_index++; }
    if(quotation){ localbuf[output_index] = '>'; output_index++; }
    return output_index;
}
```

Good:



Bad:



```

input = "Name Lastname < name@mail.org >
()()()()()()()()()()()()()()()()()()()()()()()()()()()()
()()()()()()()()()()()()()()()()()()()()()()()()()()()()
()()()"
  
```

Source: Cracking Sendmail crackaddr Still a challenge for automated program analysis?

# Activity: Does this program contain a vulnerability?

```
int main(int argc, char *argv[]) {
    int64_t x = strtoll(argv[1], NULL, 10);
    char buf[64];
    if (x <= 2 || (x & 1) != 0)
        return 1;
    int64_t i;
    for (i = x; i > 0; i--)
        if (foo(i) && foo(x - i))
            return 1;
    strcpy(buf, argv[2]); // reachable ?
}
```

```
int foo(int64_t x) {
    int64_t i, s;
    for (i = x - 1; i >= 2; i--)
        for (s = x; s >= 0; s -= i)
            if (s == 0)
                return FALSE;
    return TRUE;
}
```

# Activity: Does this program contain a vulnerability?

```
int main(int argc, char *argv[]) {
    int64_t x = strtoll(argv[1], NULL, 10);
    char buf[64];
    // x is an even number that is greater than 2
    if (x <= 2 || (x & 1) != 0)
        return 1;
    // x can be expressed as the sum of 2 primes
    int64_t i;
    for (i = x; i > 0; i--)
        if (is_prime(i) && is_prime(x - i))
            return 1;
    strcpy(buf, argv[2]); // reachable?
}
```

```
int is_prime(int64_t x) {
    int64_t i, s;
    for (i = x - 1; i >= 2; i--)
        for (s = x; s >= 0; s -= i)
            if (s == 0)
                return FALSE;
    return TRUE;
}
```

Answer: No for all 32 bit integers, but unknown for all 64bit integers...

Goldbach's conjecture:  
275+ year old unsolved math problem

# Decompositional Approach: Program Slicing

- Mark Weiser. 1981. Program slicing. In Proceedings of the 5th international conference on Software engineering (ICSE '81). IEEE Press, Piscataway, NJ, USA, 439-449.
- Ferrante, Jeanne, Karl J. Ottenstein, and Joe D. Warren. "The program dependence graph and its use in optimization." *ACM Transactions on Programming Languages and Systems (TOPLAS)* 9.3 (1987): 319-349.

# Data Flow Graph (DFG)

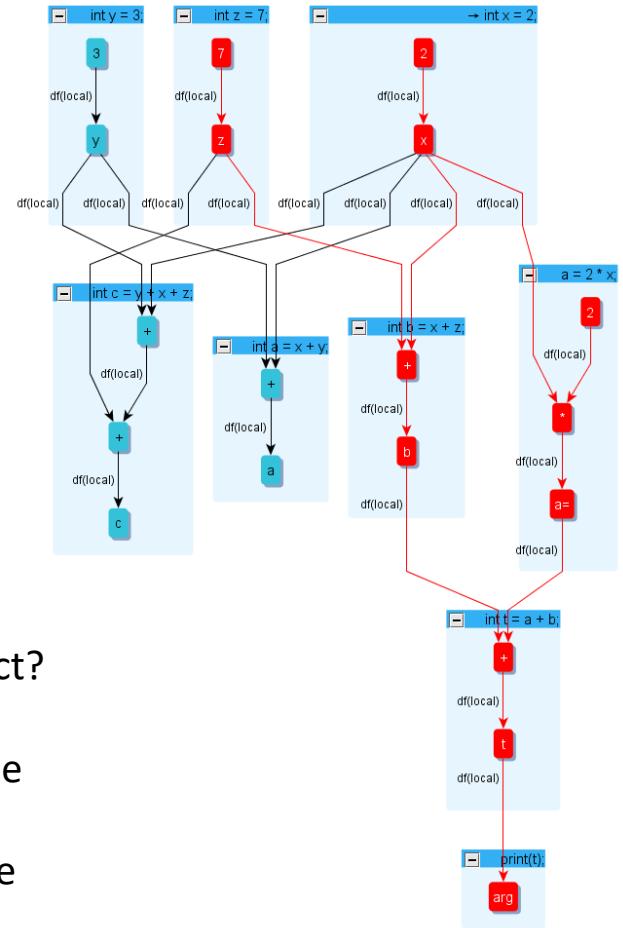
Example:

1.  $x = 2;$
2.  $y = 3;$
3.  $z = 7;$
4.  $a = x + y;$
5.  $b = x + z;$
6.  $a = 2 * x;$
7.  $c = y + x + z;$
8.  $t = a + b;$
9.  $\text{print}(t);$

Relevant lines:  
1,3,5,6,8

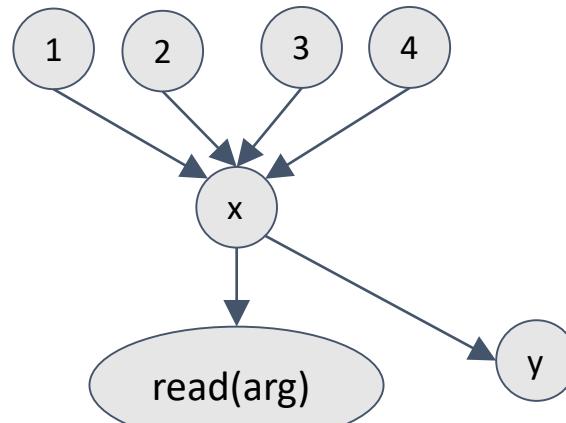
detected failure

- What lines must we consider if the value of  $t$  printed is incorrect?
- A *Data Flow Graph* creates a graph of primitives and variables where each assignment represents an edge from the RHS to the LHS of the assignment
- The *Data Flow Graph* represents global data dependence at the operator level (the atomic level) [FOW87]



## Code Transformation (before – flow insensitive): Static Single Assignment Form

1.  $x = 1;$
2.  $x = 2;$
3. if(condition)
4.    $x = 3;$
5.  $\text{read}(x);$
6.  $x = 4;$
7.  $y = x;$



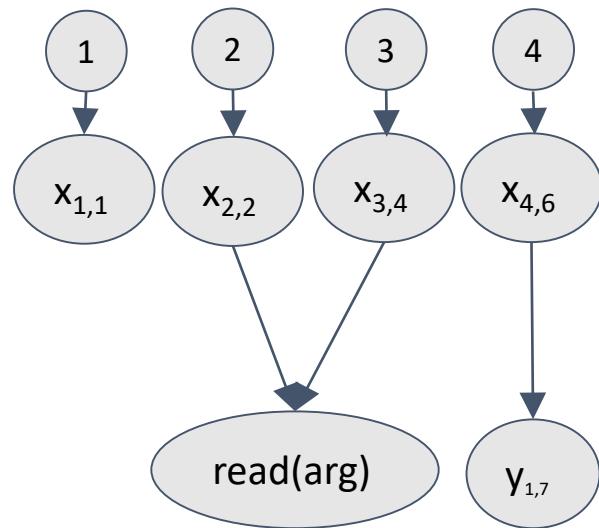
Resulting graph when statement ordering is not considered.

## Code Transformation (after – flow sensitive): Static Single Assignment Form

1.  $x = 1;$   
 2.  $x = 2;$   
 3. if(condition)  
 4.    $x = 3;$   
 5. read( $x$ );  
 6.  $x = 4;$   
 7.  $y = x;$



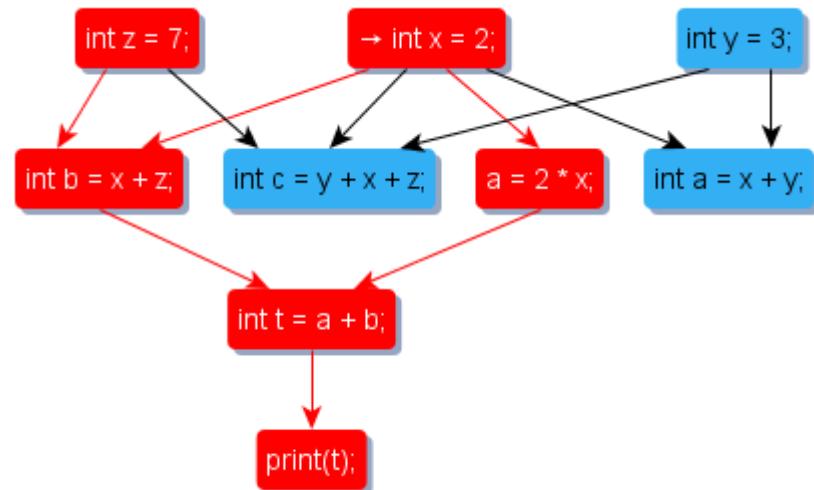
1.  $x_{1,1} = 1;$   
 2.  $x_{2,2} = 2;$   
 3. if(condition)  
 4.    $x_{3,4} = 3;$   
 5. read( $x_{2,2}, x_{3,4}$ );  
 6.  $x_{4,6} = 4;$   
 7.  $y_{1,7} = x_{4,6};$



Note: <Def#,Line#>

# Data Dependence Graph (DDG)

- Note that we could summarize data flow on a per statement level
- This graph is called a *Data Dependence Graph* (DDG)
- DDG dependences represent only the *relevant* data flow relationships of a program [FOW87]



# Data Dependence Slicing

- Reverse Data Dependence Slice
  - What statements influence the assigned value in this statement?
- Forward Data Dependence Slice
  - What statements could the assigned value in this statement influence?

# Control Flow Graph (CFG)

Example:

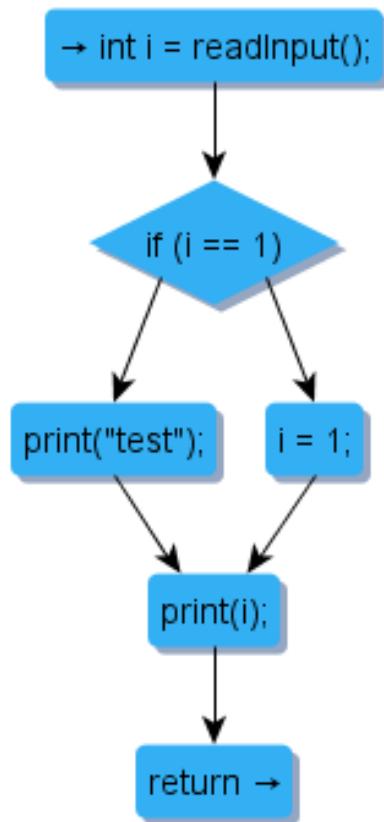
```

1. i = readInput();
2. if(i == 1)
3.     print("test");
    else
4.     i = 1;
5. print(i); ← detected failure
6. return; // terminate
  
```

Relevant lines:  
1,2,4

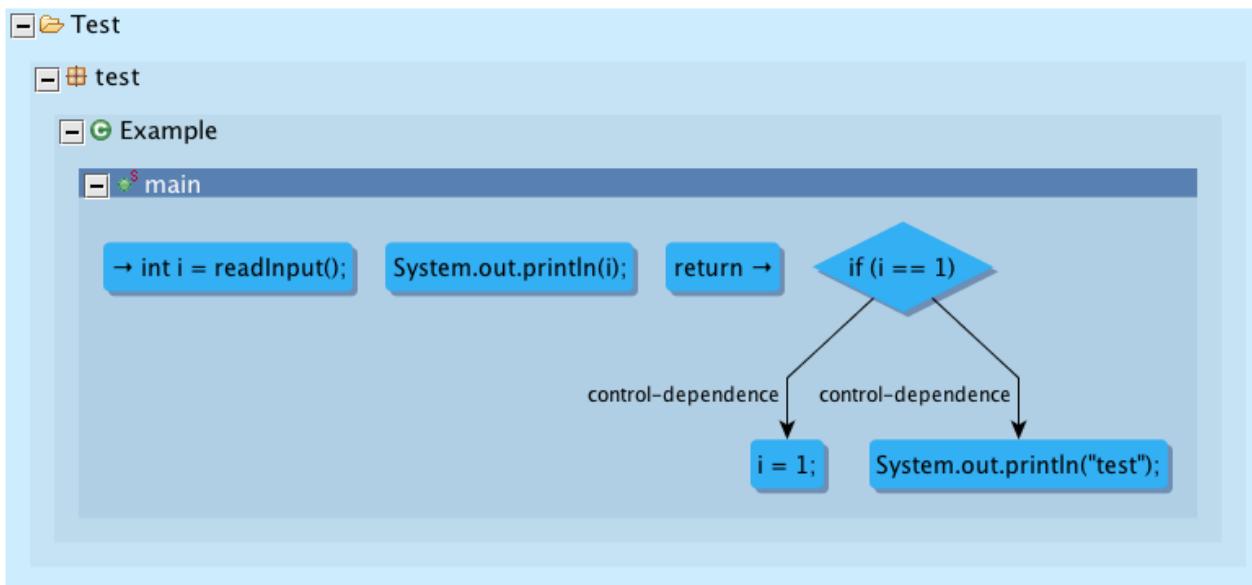
What lines must we consider if the value of *i* printed is incorrect?

- A *Control Flow Graph (CFG)* represents the possible sequential execution orderings of each statement in a program
- Data flow influences control flow, so this graph is not enough



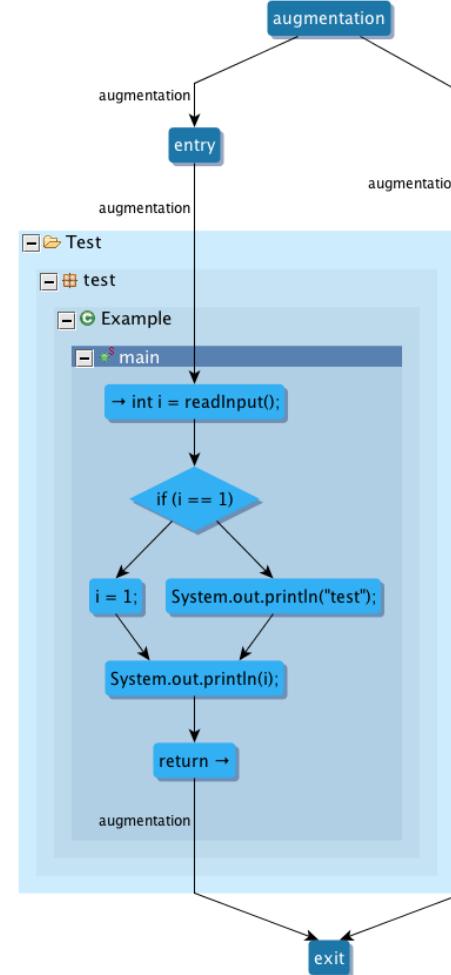
# Control Dependence Graph (CDG)

- If a statement X determines whether a statement Y can be executed then statement Y is *control dependent* on X
- Control dependence exists between two statements, if a statement directly controls the execution of the other statement [FOW87]



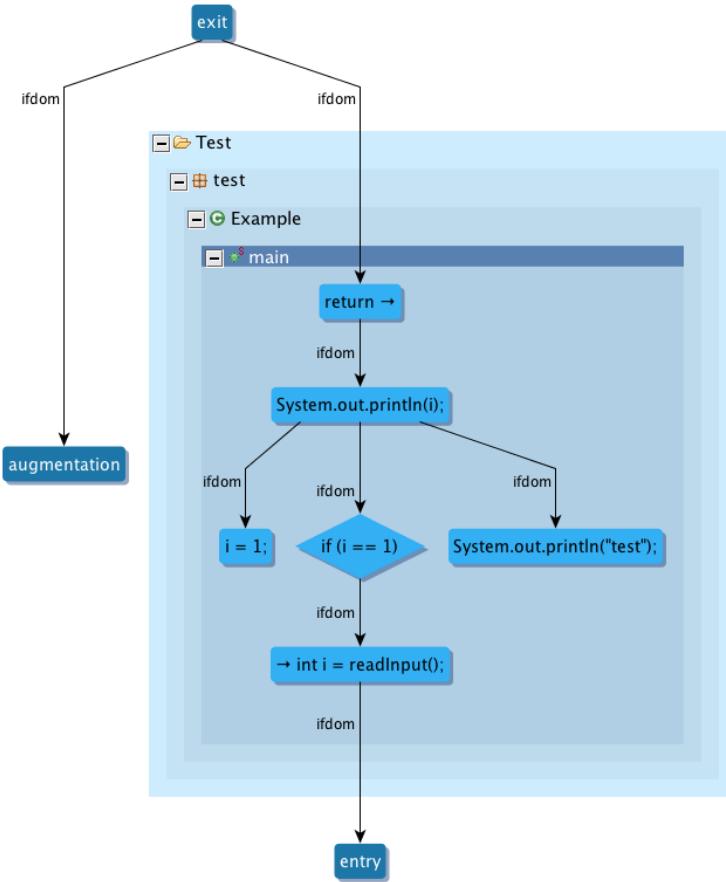
# Building a CDG (1)

- First augment the CFG with a single “entry” node and single “exit” node.
- Create an “augmentation” node which has the “entry” and “exit” nodes as children.



# Building a CDG (2)

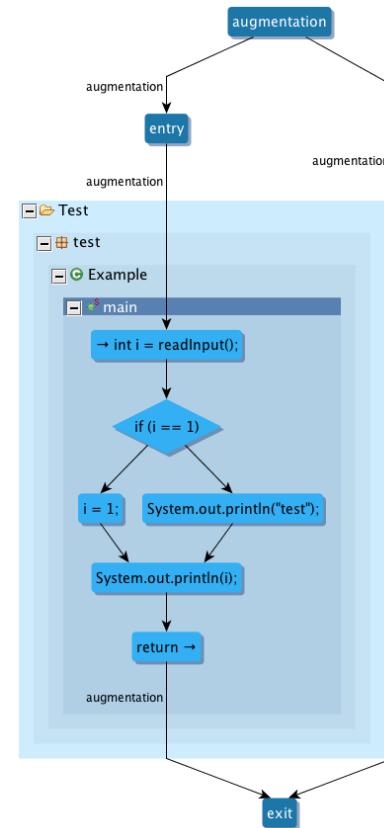
- X *dominates* Y if every path from the entry node to Y must go through X
- A *dominator tree* is a tree where each node's children are those nodes it immediately dominates
- Compute a forward dominance tree (i.e. post-dominance analysis) of the augmented CFG



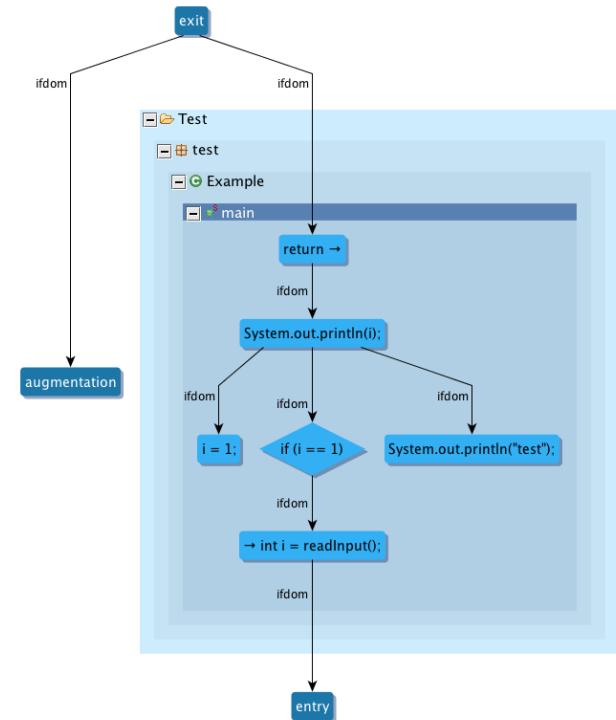
# Building a CDG (3)

- The *least common ancestor* (LCA) of two nodes X and Y is the deepest tree node that has both X and Y as descendants
- For each edge  $(X \rightarrow Y)$  in CFG, find nodes in FDT from  $\text{LCA}(X,Y)$  to Y, which are *control dependent* on X.
  - Exclude  $\text{LCA}(X,Y)$  if  $\text{LCA}(X,Y)$  is not X

Augmented CFG (ACFG)



Forward Dominance Tree (FDT)



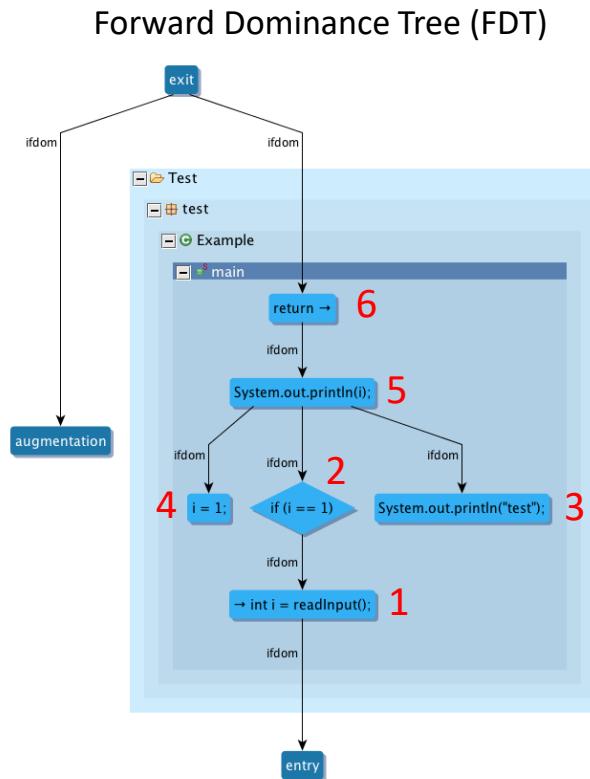
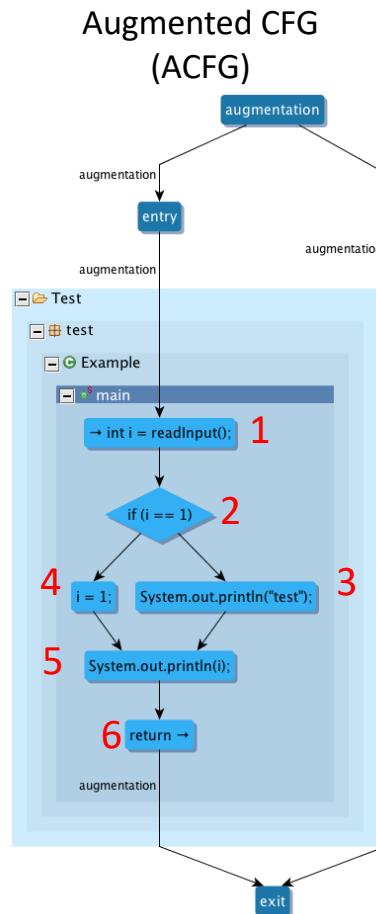
# Building a CDG (4)

Edge X→Y in ACFG	LCA(X,Y) in FDT	FDT Nodes Between(LCA, Y)
1 → 2	2	2
2 → 3	5	5, 3
2 → 4	5	5, 4
4 → 5	5	5
3 → 5	5	5
5 → 6	6	6

Note: Remove LCA(X,Y) if LCA(X,Y) != X

## Example:

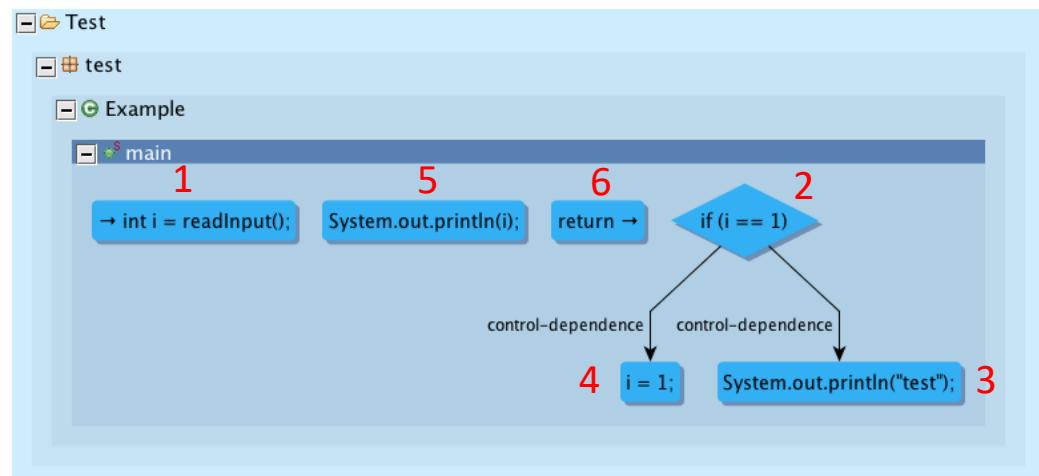
1. i = readInput();
2. if(i == 1)
3. print("test");
- else
4. i = 1;
5. print(i);
6. return; // terminate program



# Control Dependence Graph

Edge X→Y in ACFG	LCA(X,Y) in FDT	FDT Nodes Between(LCA, Y)
1 → 2	2	2
2 → 3	5	5, 3
2 → 4	5	5, 4
4 → 5	5	5
3 → 5	5	5
5 → 6	6	6

FDT Nodes Between(LCA, Y) are *Control Dependent* on X.



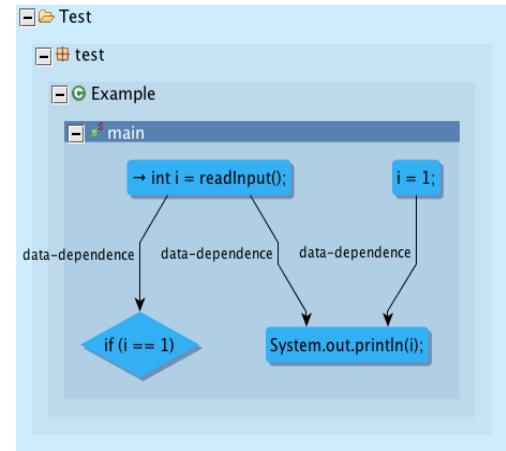
# Control Dependence Slicing

- Reverse Control Dependence Slice
  - What statements does this statement's execution depend on?
- Forward Control Dependence Slice
  - What statements could execute as a result of this statement?

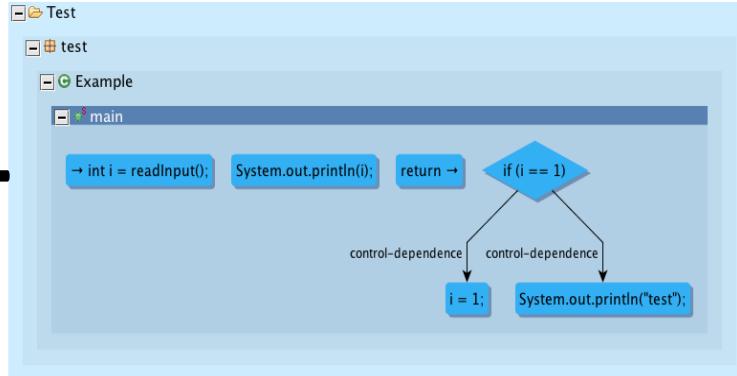
# Program Dependence Graph (PDG)

- Both DDG and CDG nodes are statements
- The union of a DDG and the CDG is a PDG

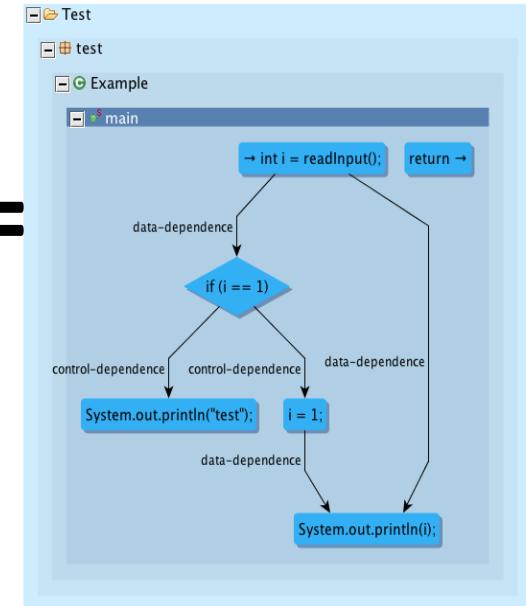
DDG



CDG



PDG



# Program Slicing (Impact Analysis)

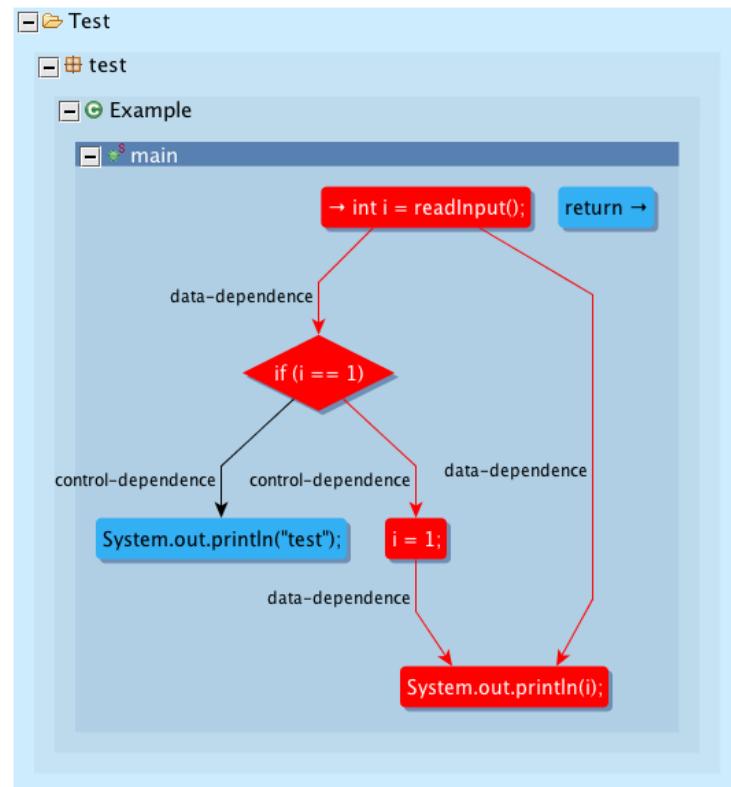
- Reverse Program Slice
  - Answers: What statements does this statement's execution depend on?
- Forward Program Slice
  - Answers: What statements could execute as a result of this statement?

Example:

```

1. i = readInput();
2. if(i == 1)
3.     print("test");
   else
4.     i = 1;
5. print(i); ← detected failure
6. return; // terminate
    
```

Relevant lines:  
1,2,4



# Taint Analysis

- How can we track the flow of data from the source ( $x$ ) to the sink ( $y$ )?
- Neither DFG/DDG nor CFG/CDG alone are enough to answer whether  $x$  flows to  $y$
- Taint = (forward slice of *source*) intersection (reverse slice of *sink*)

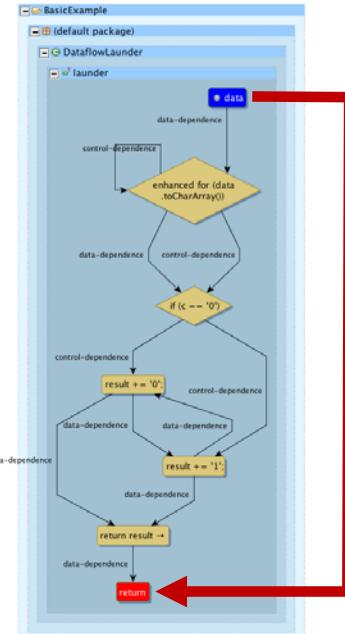
```
public class DataflowLaunder {  
  
    public static void main(String[] args) {  
        String x = "1010";  
        String y = launder(x);  
        System.out.println(y + " is a laundered version of " + x);  
    }  
  
    public static String launder(String data){  
        String result = "";  
        for(char c : data.toCharArray()){  
            if(c == '0')  
                result += '0';  
            else  
                result += '1';  
        }  
        return result;  
    }  
}
```

```
DataflowLaunder.java ✘
```

```

1 /**
2 * A toy example of laundering data through "implicit dataflow paths"
3 * The launder method uses the input data to reconstruct a new result
4 * with the same value as the original input.
5 *
6 * @author Ben Holland
7 */
8
9 public class DataflowLaunder {
10
11     public static void main(String[] args) {
12         String x = "1010";
13         String y = launder(x);
14         System.out.println(y + " is a laundered version of " + x);
15     }
16
17     public static String launder(String data){
18         String result = "";
19         for(char c : data.toCharArray()){
20             if(c == '0')
21                 result += '0';
22             else
23                 result += '1';
24         }
25         return result;
26     }
27
28 }
```

```
Taint Graph ✘
```



The Taint Graph visualization shows the data flow in the 'DataflowLaunder' program. It starts with a blue node labeled 'data' at the top. An arrow labeled 'data-dependence' points down to a yellow diamond node labeled 'enhanced for (data.toCharArray())'. From this node, two arrows labeled 'control-dependence' point to two yellow diamond nodes: 'if (c == '0')' and 'if (c == '1')'. Each of these leads to a yellow rectangle node labeled 'result += '0'' or 'result += '1''. Finally, an arrow labeled 'data-dependence' points from each of these to a yellow rectangle node labeled 'return result'. A red box highlights the entire 'launder' method node, and a red arrow points from the code editor to this highlighted area.

atlas

Problems
Javadoc
Declaration
Console
Progress
Atlas Shell (shell) ✘
Error Log
Analysis Keys
Element Detail View
Call Hierarchy

Atlas Shell (Project: shell)

```

var taint = new com.ensoftcorp.open.slice.analysis.TaintGraph(source, sink)

taint: com.ensoftcorp.open.slice.analysis.TaintGraph = com.ensoftcorp.open.slice.analysis.TaintGraph@13df7ce4

show(taint.getGraph(), taint.getHighlighter(), title="Taint Graph")
```

Evaluate: <type an expression or enter ":help" for more information>

<https://github.com/EnSoftCorp/slicing-toolbox>

## Atlas Extensible Common Software Graph (XCSG) Schema

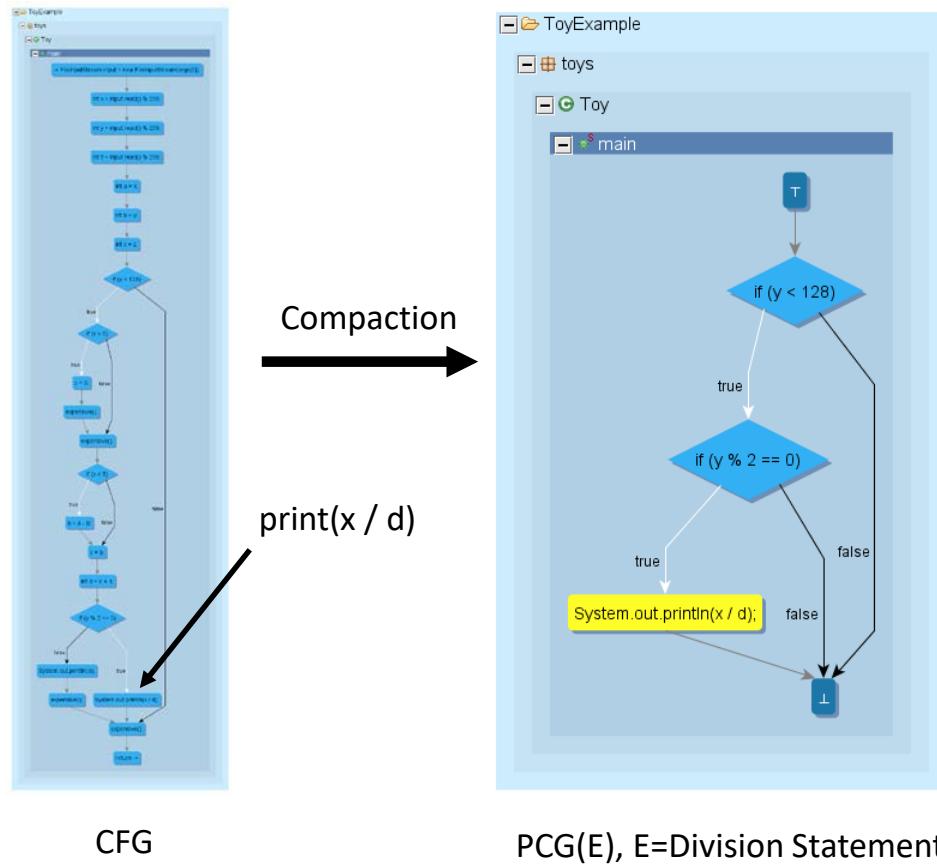
- XCSG's use of XCSG.DataFlow\_Edge and XCSG.ControlFlow\_Edge are compatible with the definitions of control and data flow as put forward by FOW87 paper
- [https://ensoftatlas.com/wiki/Extensible\\_Common\\_Software\\_Graph](https://ensoftatlas.com/wiki/Extensible_Common_Software_Graph)

# Slicing Continuations

- System Dependence Graphs
  - Reps, Thomas, Susan Horwitz, and Mooly Sagiv. "Precise interprocedural dataflow analysis via graph reachability." *Proceedings of the 22nd ACM SIGPLAN-SIGACT symposium on Principles of programming languages*. ACM, 1995.
- Survey of Slicing Techniques
  - Tip, Frank. *A survey of program slicing techniques*. Centrum voor Wiskunde en Informatica, 1994.

# Projected Control Graphs

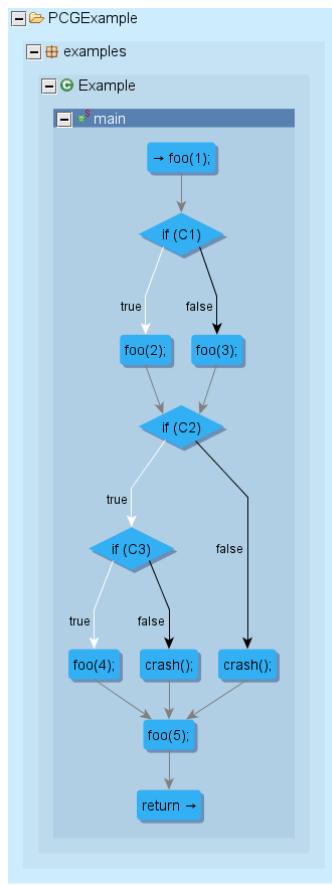
- Projected Control Graph (PCG)
  - Proposed by Ahmed Tamrawi in 2016
  - Efficiently groups program behaviors into equivalence classes of homomorphic behaviors
    - Parameterized by events of interest
    - Only event statements and necessary conditions are retained
  - Result is a compact structure-preserving model of a CFG w.r.t. events of interest



# Concept: Projected Control Graph

```

1 public void main() {
2     foo(1);
3     if (C1) {
4         foo(2);
5     } else {
6         foo(3);
7     }
8     if (C2) {
9         if (C3) {
10            foo(4);
11        } else {
12            crash();
13        }
14    } else {
15        crash();
16    }
17    foo(5);
18    return;
19 }
```



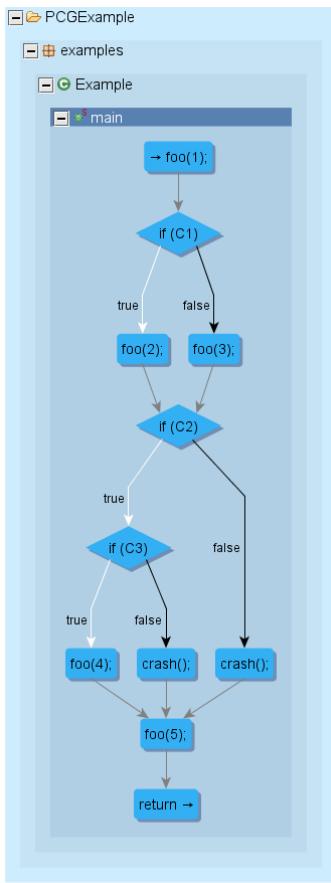
$2^3=8$  possible values for the tuple (C1, C2, C3)

C1	C2	C3	Behavior
False	False	False	
False	False	True	
False	True	False	
False	True	True	
True	False	False	
True	False	True	
True	True	False	
True	True	True	

# Concept: Projected Control Graph

```

1 public void main() {
2     foo(1);
3     if (C1) {
4         foo(2);
5     } else {
6         foo(3);
7     }
8     if (C2) {
9         if (C3) {
10            foo(4);
11        } else {
12            crash();
13        }
14    } else {
15        crash();
16    }
17    foo(5);
18    return;
19 }
```



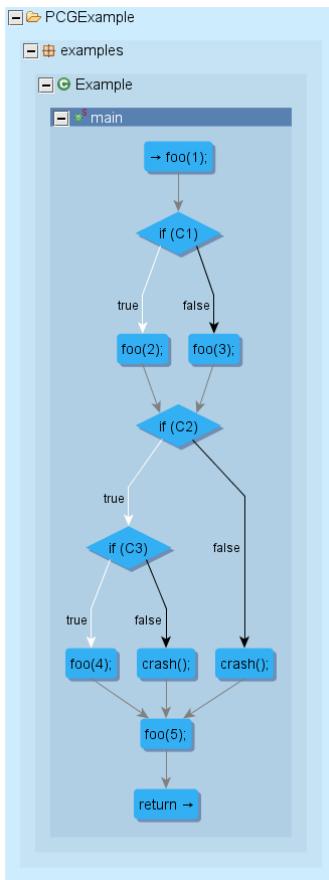
If C2 is false then C3 is not evaluated.

C1	C2	C3	Behavior
False	False	N/A	
False	False	N/A	
False	True	False	
False	True	True	
True	False	N/A	
True	False	N/A	
True	True	False	
True	True	True	

# Concept: Projected Control Graph

```

1 public void main() {
2     foo(1);
3     if (C1) {
4         foo(2);
5     } else {
6         foo(3);
7     }
8     if (C2) {
9         if (C3) {
10            foo(4);
11        } else {
12            crash();
13        }
14    } else {
15        crash();
16    }
17    foo(5);
18    return;
19 }
```



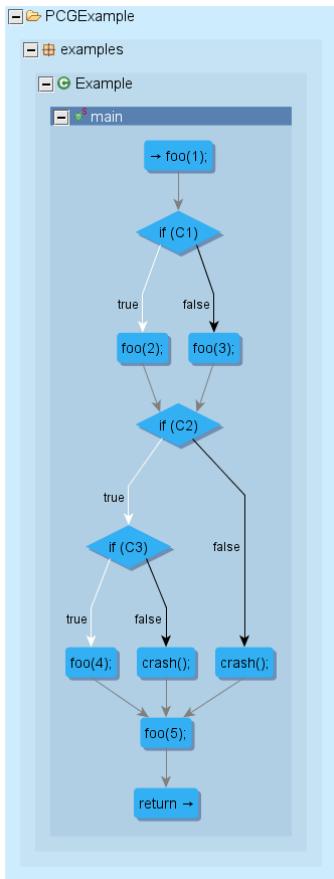
CFG has 6 paths.

C1	C2	C3	Behavior
False	False	N/A	
False	True	False	
False	True	True	
True	False	N/A	
True	True	False	
True	True	True	

# Concept: Projected Control Graph

```

1 public void main() {
2     foo(1);
3     if (C1) {
4         foo(2);
5     } else {
6         foo(3);
7     }
8     if (C2) {
9         if (C3) {
10            foo(4);
11        } else {
12            crash();
13        }
14    } else {
15        crash();
16    }
17    foo(5);
18    return;
19 }
```



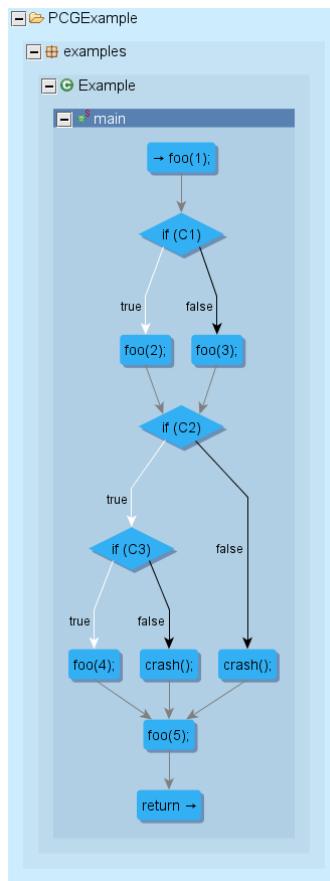
What paths include a “crash” event?

C1	C2	C3	Behavior
False	False	N/A	2,3,6,8,15,17,18
False	True	False	2,3,6,8,9,12,17,18
False	True	True	2,3,6,8,9,10,17,18
True	False	N/A	2,3,4,8,15,17,18
True	True	False	2,3,4,8,9,12,17,18
True	True	True	2,3,4,8,9,10,17,18

# Concept: Projected Control Graph

```

1 public void main() {
2     foo(1);
3     if (C1) {
4         foo(2);
5     } else {
6         foo(3);
7     }
8     if (C2) {
9         if (C3) {
10            foo(4);
11        } else {
12            crash();
13        }
14    } else {
15        crash();
16    }
17    foo(5);
18    return;
19 }
```



4 of 6 behaviors have “crash” events.

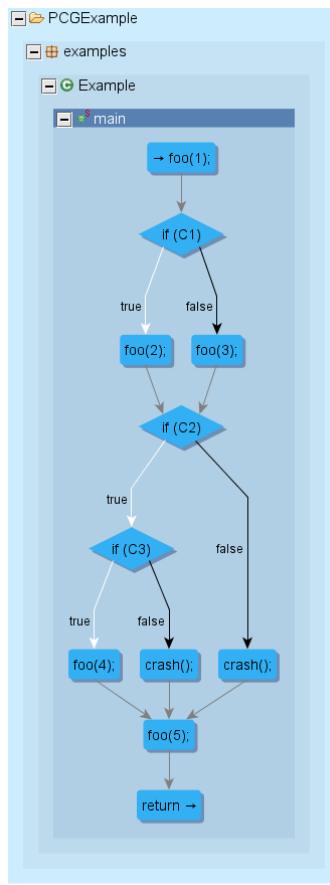
2 of 6 behaviors do not have “crash” events.

C1	C2	C3	Behavior
False	False	N/A	2,3,6,8, <b>15,17,18</b>
False	True	False	2,3,6,8,9, <b>12,17,18</b>
False	True	True	<del>2,3,6,8,9,10,17,18</del>
True	False	N/A	2,3,4,8, <b>15,17,18</b>
True	True	False	2,3,4,8,9, <b>12,17,18</b>
True	True	True	<del>2,3,4,8,9,10,17,18</del>

# Concept: Projected Control Graph

```

1 public void main() {
2     foo(1);
3     if (C1) {
4         foo(2);
5     } else {
6         foo(3);
7     }
8     if (C2) {
9         if (C3) {
10            foo(4);
11        } else {
12            crash();
13        }
14    } else {
15        crash();
16    }
17    foo(5);
18    return;
19 }
```



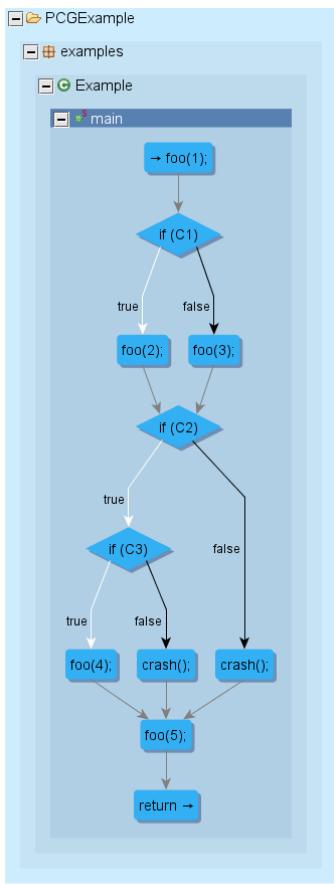
There are 3 unique behaviors w.r.t. “crash” events.

C1	C2	C3	Homomorphic Behavior
False	False	N/A	8,15
False	True	False	8,9,12
False	True	True	8,9
True	False	N/A	8,15
True	True	False	8,9,12
True	True	True	8,9

# Concept: Projected Control Graph

```

1 public void main() {
2     foo(1);
3     if (C1) {
4         foo(2);
5     } else {
6         foo(3);
7     }
8     if (C2) {
9         if (C3) {
10            foo(4);
11        } else {
12            crash();
13        }
14    } else {
15        crash();
16    }
17    foo(5);
18    return;
19 }
```



2 of 3 homomorphic behaviors (B1, B2) have “crash” events.

1 of 3 homomorphic behaviors (B3) are void w.r.t. “crash” events.

C1	C2	C3	Homomorphic Behavior
False	False	N/A	8,15
False	True	False	8,9,12
False	True	True	8,9
True	False	N/A	8,15
True	True	False	8,9,12
True	True	True	8,9

# PCGs in Practice

- We research tools and techniques that:
  - Produce human-verifiable and comprehensible evidence
  - Are scalable
- Verification of lock/unlock pairs in Linux
  - Over 66,000 instances (3 versions of Linux)
  - L-SAP tool has resulted in 8 bug reports which have been fixed
- Resulted in key concept: Projected Control Graph (PCG)
  - Retains minimal necessary program behaviors

# Motivation Behind PCGs

- Control Flow Graphs can be HUGE
  - CFG of a single function in Linux
    - lustre\_assert\_wire\_constants function has  $2^{656}$  paths!
  - Programmers don't play dice...
  - Is all of this complexity necessary to solve a given analysis problem?
  - Is there an underlying design pattern that can be leveraged to complete the analysis?





Military Communications for the 21st Century  
October 29-31, 2018 • LAX Marriott, Los Angeles, CA

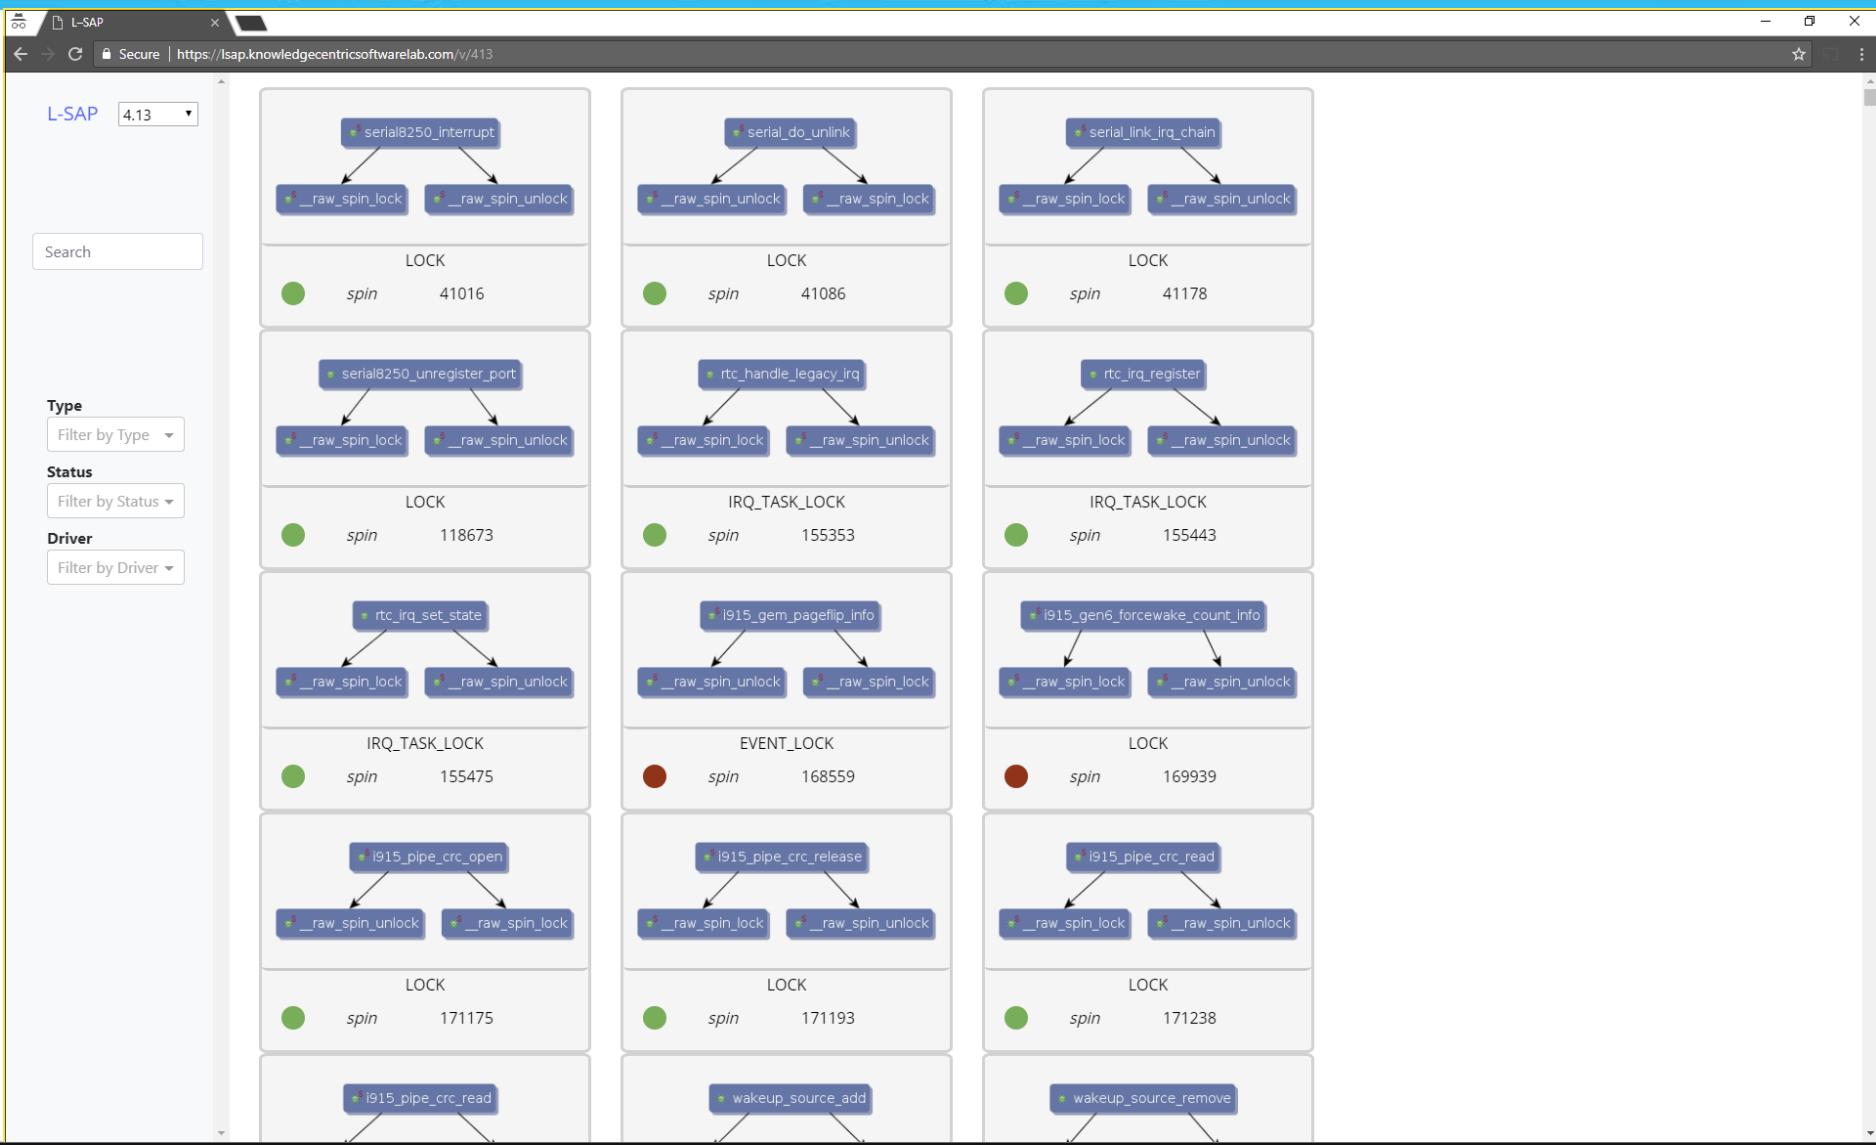


# PCG Example from Linux Kernel

- Example usage: verifying Linux lock/unlock pairs. Lock must be followed by unlock on all paths.
  - Research group found 8 unknown and now accepted bugs in Linux 3.13 to 4.3, 13 reported
  - Real Example from Linux 3.13: Red=LOCK, Green=UNLOCK; this PCG reveals the CFG has a path (via a missing switch default statement) where a LOCK can occur without a corresponding UNLOCK.



<https://github.com/EnSoftCorp/pcg-toolbox>  
<https://kcsli.github.io/L-SAP/bugs> (Linux study results)



L-SAP 4.13

Secure | https://lsap.knowledgecentricsoftwarelab.com/v/413/i/246193

drivers/tty/serial/8250/8250\_port.c

Locking type: spin  
Instance ID: 246193  
Length: 38  
Title: lock  
Offset: 58110  
Status: PAIRED

**Old Version** **New Version**

**Type**  
Filter by Type ▾

**Status**  
Filter by Status ▾

**Driver**  
Filter by Driver ▾

`serial8250_do_startup`

```
graph TD; serial8250_do_startup --> raw_spin_unlock; serial8250_do_startup --> raw_spin_lock
```

**Control Flow Graphs**



**Projected Control Graphs**



# CFG to PCG Reduction Statistics

- For Linux 3.19-rc1, the average CFG to PCG reduction is 73.4% for the number of nodes and 75.5% for the number of edges
- Extreme example from Linux study: client\_common\_fill\_super
  - CFG: 1, 101 nodes, 1, 179 edges, and 249 branch nodes
  - PCG: 15 nodes, 28 edges, and 13 branch nodes

# Comparison to BLAST

- Berkeley Lazy Abstraction Software Verification Tool (BLAST) is a top-rated tool in the software verification competition (SV-COMP).
- BLAST verifies 43,766 (65.7)% of Lock instances as safe, and it is inconclusive (crashes or times out) on 22,843 instances. BLAST does not find any unsafe instances.
- The PCG-based automated verification tool verifies 66,151 (99.3)% of Lock instances as safe, and it is inconclusive on 451 instances. One of the bugs we found is an instance that BLAST reported as a safe instance.

Discovered Bugs – L-SAP ×

Secure | https://knowledgecentricsoftwarelab.com/L-SAP/bugs/

L-SAP  
Scalable and Accurate Linux Lock/Unlock Pairing

Home Install Tutorials Bugs

## Discovered Bugs

Below is a list of all bug cases that have been discovered with L-SAP.

### Bug - 10/18/2015

**Version:** 4.3-rc6

**Source File:** /v4.3-rc6/drivers/media/pci/ddbridge/ddbridge-core.c

Function `tuner_attach_tda18271` (line 595) acquires a lock on `port->i2c_gate_lock` via the call to the function pointer (`input->fe->ops.i2c_gate_ctrl(input->fe, 1)`) (line 601) which calls function `drxk_gate_ctrl`. However, when the call to function `dvb_attach` on line 602 fails (returns `NULL`), the lock on `port->i2c_gate_lock` is never released.

The bug was currently on the master branch of v4.3-rc6 as of writing the paper.

### Bug - Paper Case Study I

**Version:** 3.19-rc1

**Source File:** /v3.19-rc1/drivers/mmc/host/toshsd.c

Function `toshsd_thread_irq` has an unpaired lock `spin_lock_irqsave(&host->lock, flags)` at line 176. The bug occurs when the function returns at line 179 without unlocking the spin object `host->lock`.

The bug was spread across all the release candidates of 3.19-rc1 through 3.19-rc7 and got fixed in 4.0. The fixing commit ID is: 8a66fdae771487762519db0546e9ccb648a2f911.

### Bug - Paper Case Study II

**Version:** 3.19-rc1

**Source File:** /v3.19-rc1/drivers/scsi/fnic/fnic\_fcs.c

Function `fnic_handle_fip_timer` has an unpaired lock

# Interactive Audit of Android Application

- ConnectBotBad
  - Several thousand lines of source code
  - Has multiple malwares
  - Work smarter not harder
  - Use control flow, data flow, program slices to test hypotheses
  - Leverage some knowledge of Android APIs to search sensitive interactions
- FlashBang
  - Example from the wild (but decompiled and refactored for this lab)
  - Try auditing the source code version
  - Try auditing the binary version
- Timelapse Audit of DARPA APAC Challenge Application
  - <https://www.youtube.com/watch?v=p2mhfOMmgKI>

# Additional Resources

- <https://www.ece.iastate.edu/kcls>
- Reach out to us and start a discussion!
  - Suresh Kothari – [kothari@iastate.edu](mailto:kothari@iastate.edu)
  - Benjamin Holland – [bholland@iastate.edu](mailto:bholland@iastate.edu)

IOWA STATE UNIVERSITY  
Knowledge-Centric Software Lab

Search 

**People**

- [People](#)
- [Tools](#)
- [Atlas](#)
- [Atlas Toolboxes](#)
- [FlowMiner](#)
- [L-SAP](#)

**Publications**

- Book Chapters (1)
- Competitions (1)
- Papers (29)
- Short Courses (3)
- Talks (9)
- Tutorials (11)
- Upcoming (3)

**Monthly Activity**

- July 2018 (4)
- June 2018 (1)
- May 2018 (1)
- April 2018 (1)
- March 2018 (1)
- December 2017 (3)
- October 2017 (2)
- September 2017 (1)
- July 2017 (2)
- June 2017 (1)
- March 2017 (3)
- December 2016 (1)
- November 2016 (2)
- October 2016 (4)
- September 2016 (2)
- August 2016 (1)
- May 2016 (4)
- December 2015 (2)
- November 2015 (2)
- October 2015 (1)
- May 2015 (1)
- December 2014 (2)
- October 2014 (1)
- September 2014 (1)
- May 2014 (1)

**People**

The Knowledge-Centric Software Laboratory at the Department of Electrical and Computer Engineering consists of a group of researchers interested in solving hard problems in software program analysis. The laboratory is led and directed by professor and entrepreneur Dr. Suresh Kothari.

Recent research funding has come primarily from DARPA contracts FA8750-12-2-0126 and FA8750-15-2-0080.

**Director**

- [Suresh Kothari](#) (Richardson Professor)

**Current Members**

- [Benjamin Holland](#) (PhD Student)
- [Derrick Lockwood](#) (Masters Student)
- [Le Zhang](#) (PhD Student)
- [Payas Avadhutkar](#) (PhD Student)
- [Shanwan Ram](#) (Visiting Scholar)

**Past Members**

Ahmed Tamrawi, Akshay Deepak, Aravind Krishnamoorthy, Curtis Utterich, Daman Singh, Dan Harvey, Dan Stine, Ganesh Ram Santhanam, Jaekyu Cho, Jason Stanek, Jeremias Sauceda, Jim Carl, Jon Matthews, Kang Gyu, Lisbeth Meum, Luke Bishop, Murali Ravirala, Nithil Ranade, Sandeep Krishnan, Sergio Ferrero, Simanta Mitra, Srinivas Neginalal, Tom Deering, Wei Ke, Xiaozheng Ma, Yogy Namrata, Youngtae Kim, Yunbo Deng, Zach Lones

